Autonomy: Executive Summary

> **INTERNAL DOCUMENT** — Contains implementation details, file paths, and architecture specifications.

What Your Supply Chain Organization Looks Like After Autonomy

Version: 4.1 **Date**: March 9, 2026

Executive Overview

Most supply chain organizations operate the same way they did a decade ago. Planners open spreadsheets or legacy screens every Monday, review thousands of SKU-level exceptions, make judgment calls under time pressure, and hope the plan holds until next week. When disruptions hit mid-cycle, the response is reactive — phone calls, expedited shipments, and costly overtime. The institutional knowledge that makes this work lives in the heads of a few senior planners, and when they leave, it leaves with them.

Autonomy changes the operating model. Not by replacing your planning team, but by restructuring what they spend their time on — shifting the workforce from routine plan generation to exception governance, from reactive firefighting to proactive risk management, and from tribal knowledge to institutional learning that compounds over time.

This document describes what is different — for the organization, for daily operations, and for the people doing the work — when Autonomy is fully adopted.

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Agent Architecture

```
graph TD
    SOP["<b>Strategic Consensus</b><br/><i>Weekly · Policy Parameters</i>"]
    AAP["<b>Cross-Authority Protocol</b><br/><i>Ad Hoc · Trade-off Authorization</i>"]
    NET["<b>Network Coordination</b><br/><i>Daily · Inter-Site Directives</i>"]
    SITE["<b>Site Cross-Agent Coordination</b><br/><i>Hourly · Predictive Urgency Modulation

    subgraph EXEC["<b>11 Execution Agents</b> · &lt;10ms per Decision"]
        direction LR
        DEMAND["Demand Sensing<br/><small>Order Promising · Exception Tracking</small>"]
        SUPPLY["Supply Securing<br/><small>Purchasing · Transfers · Subcontracting</small>"]
        HEALTH["Health Monitoring<br/><small>Inventory Buffers · Forecast Adjustment</small>
        INTEG["Integrity Protection<br/><small>Quality · Maintenance</small>"]
        BUILD["Production & Logistics<br/><small>Manufacturing Orders · Transfer Orders</sma
    end

    SOP -->|"policy parameters"| NET
    AAP <-->|"authorization"| NET
    NET -->|"site directives"| SITE
    SITE -->|"urgency adjustments"| EXEC
    EXEC -->|"signals & outcomes"| SITE
    SITE -->|"escalation"| NET
```

Part 1: What Changes for the Organization

From Periodic Planning to Continuous Response

Today, most organizations plan on a fixed cadence. Weekly MPS runs. Monthly S&OP meetings. Quarterly network reviews. The problem is that supply chains don't wait for cadence. A supplier delay on Tuesday doesn't care that your MPS run was Monday.

With Autonomy, the planning system is always on. When an event occurs — a supplier misses a delivery window, a customer order spikes unexpectedly, a production line goes down — the system detects it, evaluates the downstream impact, generates a recommended response, and either executes within pre-approved guardrails or surfaces the decision to the right person with a fully evaluated set of options.

The shift is structural:

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

	Before Autonomy	After Autonomy
Planning trigger	Calendar (weekly/monthly)	Events (as they happen)
Scope of each plan run	Entire portfolio, all SKUs	Only affected products and locations
Response latency	Days to a week	Minutes to hours
Human role	Generate the plan	Govern the plan
Knowledge capture	Undocumented, in people's heads	Every decision and its reasoning recorded
Improvement mechanism	Annual consultant engagement	Continuous, automatic from daily work

From Cost Center to Competitive Advantage

In most organizations, the supply chain planning function is seen as a cost of doing business — necessary, expensive, and difficult to improve incrementally. Autonomy repositions it.

When the planning system learns from every decision, every override, and every outcome, it creates a **compounding knowledge asset**. The more decisions your team makes, the smarter the system gets. The smarter the system gets, the more decisions it can handle without human intervention. The more it handles autonomously, the more time your team has for strategic work — network redesign, supplier negotiations, new product introduction planning — that actually moves the business forward.

This compounding loop is difficult for competitors to replicate because it is built on your data, your team's judgment, and your specific operating context.

From Fragmented Tools to One Operating System

A typical mid-market supply chain team juggles between their ERP (for transactions), a planning tool or spreadsheets (for MPS/MRP), email and phone (for collaboration), separate BI dashboards (for analytics), and manual processes (for exception handling). Each tool has its own data, its own logic, and its own version of the truth.

Autonomy consolidates the planning workflow into a single environment where demand signals, supply plans, inventory positions, order status, exception alerts, and AI recommendations are all visible in context. Planners don't context-switch between systems. They work in one place, and every action — whether human or AI — is tracked, explained, and available for review.

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Part 2: What Changes for Daily Operations

The Daily Workflow, Redesigned

Monday morning, before Autonomy:

A planner arrives, opens the planning system, and begins reviewing the weekly MPS output. There are 847 exceptions flagged. She knows from experience that most of them are noise — parameter mismatches, minor demand shifts that will self-correct, safety stock recalculations that don't need action. But she can't tell which of the 847 actually matter without clicking into each one. By Thursday, she's reviewed 400. The other 447 roll into next week. Three of them were urgent. One caused a customer stockout.

Monday morning, with Autonomy:

The same planner arrives to a prioritized worklist of 14 items. Overnight, the system processed the same 847 potential exceptions. It auto-resolved 780 within pre-approved guardrails (reorder quantities adjusted, safety stock recalculated, minor schedule shifts executed). It flagged 53 for informational review — actions it took that the planner should be aware of but doesn't need to approve. And it escalated 14 that require human judgment: a supplier bankruptcy risk, a demand spike for a promotional item that exceeds capacity, a quality hold on a key component.

For each of those 14 items, she sees: - What happened (the triggering event) - What the system recommends (ranked options with trade-offs) - Why it recommends it (reasoning grounded in specific data — this order, this inventory level, this lead time) - What happens if she does nothing (projected impact on service level, cost, and downstream orders)

She resolves all 14 by 10 AM. The three urgent items from the old workflow? They were auto-resolved Friday evening.

How Decisions Flow

The decision architecture follows a principle we call **Automate-Inform-Inspect-Override**:

Automate: Routine decisions execute within guardrails without human involvement. Guardrails are business rules set by the planning team: maximum order value, maximum safety stock change, minimum service level floor, cost increase ceiling. Anything within bounds executes automatically.

Inform: When the system takes an action, it logs the decision and notifies the relevant planner. Not as a pop-up or an urgent alert — as a daily digest or a section of the worklist marked "for your awareness." The planner can review at their pace.

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Inspect: For decisions that exceed guardrails or involve unusual circumstances, the system presents a structured recommendation. The planner can drill into the reasoning, examine the underlying data, ask follow-up questions in plain language, and compare alternative scenarios side by side.

Override: When a planner disagrees with a recommendation, they don't just click "reject." They provide context — why they're overriding, what they know that the system doesn't, what external factor is at play. This override, with its reasoning, becomes training data. The system learns not just that the planner disagreed, but why, and adjusts its future recommendations accordingly.

Over time, the override rate drops — not because the system forces compliance, but because it gets better at incorporating the judgment patterns that drive human overrides in the first place.

Measured progression: - Week 1: ~45% of decisions auto-execute, ~35% overridden - Week 12: ~72% auto-execute, ~15% overridden - Steady state target: ~85% auto-execute, <10% overridden

Testing Changes Before They Go Live

One of the most expensive mistakes in supply chain management is implementing a policy change — a new safety stock level, a new supplier, a network redesign — and discovering weeks later that it created problems elsewhere in the network.

Autonomy includes a **digital twin** of your supply chain: a complete replica that uses the same planning logic, the same AI agents, and the same cost calculations as production. The only differences are time (it runs in fast-forward) and demand (it can use synthetic patterns or historical actuals).

This means you can: - **Test a policy change** (e.g., "What happens if we reduce safety stock by 20% at the East Coast DC?") and see the downstream impact on service levels, costs, and order patterns within hours - **Model a network change** (e.g., "What if we add a regional warehouse in Dallas?") and run hundreds of demand scenarios to measure the financial and operational impact - **Compare strategies** (e.g., "How does our current ordering approach perform against an optimized alternative across 500 simulated weeks?") and make the decision with statistical confidence instead of intuition - **Validate before deploying** any change to the production system, eliminating the "try it and hope" approach

Exception-Driven, Not Batch-Driven

Traditional planning systems process everything in batch. Every SKU, every location, every time period — whether anything changed or not. This is computationally wasteful and creates a firehose of exceptions that buries the important ones.

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Autonomy is event-driven. When something changes — a new order, a late shipment, a forecast update, a capacity constraint — only the affected products and locations are replanned. The rest of the plan holds. This means:

- **Faster response**: Critical disruptions are addressed in minutes, not days
- **Less noise**: Planners see only what actually changed, not a full regeneration of the plan
- **Better auditability**: Every plan change is tied to a specific triggering event, making root cause analysis straightforward
- **Lower computational cost**: Incremental replanning uses a fraction of the resources of full batch runs

Collaboration That Actually Works

Supply chain planning is inherently cross-functional. A procurement decision affects manufacturing capacity. A production schedule change affects distribution. A customer priority override affects allocation across the entire network.

In most organizations, this coordination happens through meetings, emails, and phone calls — slow, unstructured, and impossible to audit. Autonomy makes cross-functional coordination part of the workflow:

- **Approval workflows** route decisions to the right stakeholders automatically
- **Inline comments** on specific orders, plans, and recommendations create a threaded discussion trail tied to the artifact being discussed
- **Activity feeds** show everyone involved what changed, who changed it, and why
- **Plain-language explanations** make AI recommendations accessible to non-technical stakeholders — a VP of Supply Chain can understand why the system recommends a particular action without needing a data science background

---

Part 3: What Changes for the Workforce

The Planner's Role Evolves

The most significant workforce impact is what planners spend their time on. Today, an experienced planner might spend 70% of their time on routine plan generation and exception review, and 30% on strategic analysis. With Autonomy, those ratios invert.

Before: - 40% — Running and reviewing MPS/MRP outputs - 30% — Manually resolving routine exceptions - 15% — Firefighting unexpected disruptions - 10% — Strategic analysis (network design, policy optimization) - 5% — Training junior planners

After: - 5% — Reviewing the daily digest of auto-executed actions - 15% — Resolving the escalated exceptions that require human judgment - 10% — Monitoring system

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

performance and adjusting guardrails - 40% — Strategic analysis, scenario testing, policy optimization - 15% — Cross-functional collaboration (S&OP, supplier development, new product planning) - 15% — Training, override context capture, and knowledge building

This is not a reduction in headcount — it is a redeployment of expertise. The same team handles more volume, with higher quality, and spends their intellectual energy on work that differentiates the business rather than work that merely keeps it running.

Institutional Knowledge Becomes a System Asset

In a traditional planning organization, a senior planner's departure is a crisis. They carry years of context — which suppliers are reliable, which forecast adjustments are noise, which customers tolerate partial shipments, which production lines have hidden capacity. When they leave, this knowledge walks out the door.

Autonomy captures this knowledge systematically. Every override a planner makes is recorded with context. Every decision pattern is learned. Every "I adjusted this because I know the supplier is slow in Q4" becomes a data point that the system internalizes. Over months and years, this creates a **living knowledge base** — not a static document that nobody reads, but an active decision-support system that embodies the accumulated judgment of every planner who has ever worked in it.

The result: onboarding a new planner drops from 6-12 months to 2-3 months. Not because the job is simpler, but because the new planner has access to the institutional judgment that previously lived only in their predecessor's head.

Building Confidence Through Proof, Not Promises

The biggest barrier to AI adoption in supply chain is not technology — it's trust. Planners who have spent years developing their judgment are understandably skeptical when told that a system can do their job better. And they're right to be skeptical: plenty of "AI-powered" tools have overpromised and underdelivered.

Autonomy addresses this by making AI performance **observable and provable** before it touches real inventory:

1. **Side-by-side comparison**: Run the same planning scenario with human decisions and AI decisions. Compare the results — total cost, service level, inventory turns, order stability. The numbers speak.

2. **Gradual handoff**: Start with AI as a suggestion engine (copilot mode). Planners see every recommendation, accept or reject with reasoning, and maintain full control. As confidence builds, expand the guardrails. Let the system auto-execute more categories of decisions.

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

3. **Transparent reasoning**: Every AI recommendation comes with an explanation grounded in specific data — not "the algorithm says 42 units" but "given current inventory of 18 at the Dallas DC, incoming shipments of 25 arriving Thursday, and forecasted demand of 35 next week, I recommend ordering 42 units to maintain the 95% service level target while accounting for the supplier's recent 2-day lead time variability."

4. **Competitive simulation**: Planners can compete directly against the AI in simulated scenarios. Win or lose, they see exactly how the AI makes decisions, where it outperforms human intuition, and where human judgment adds value that the AI misses. This transforms skeptics into informed adopters.

The adoption timeline with this approach is typically 2-3 weeks to initial confidence, vs. 6-12 months with traditional planning software where users are simply told to trust the output.

A New Skill Set, Not a Lost One

Autonomy doesn't make supply chain expertise obsolete — it redirects it. The planning team needs to develop new competencies:

- **Guardrail design**: Setting the boundaries within which AI operates autonomously. This requires deep supply chain knowledge — you have to understand what's safe to automate and what isn't.
- **Exception governance**: Judging the 10-15% of decisions that the system escalates. These are the hard calls — the ones that require context, creativity, and cross-functional negotiation.
- **Scenario design**: Constructing meaningful tests in the digital twin. "What should we test?" is a strategic question that requires business acumen, not technical skills.
- **Override storytelling**: Providing high-quality context when overriding AI recommendations. The richer the context, the faster the system learns. This turns tacit knowledge into explicit knowledge.
- **Performance monitoring**: Tracking system accuracy, identifying drift, adjusting parameters. Think of it as tuning an instrument rather than playing every note.

These are higher-value skills. They build on existing supply chain expertise while adding a layer of analytical rigor and systems thinking that makes the team more effective — and more marketable.

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Part 4: The Business Case

What It Costs

Traditional enterprise planning (Kinaxis, SAP IBP, OMP): - Year 1: $6M+ (licenses, consulting, training, integration) - Ongoing: $1.35M/year (maintenance, consultants, user training) - Time to value: 12-18 months before operational benefits - 5-year TCO: $11.4M+

Autonomy: - Year 1: $200K (platform, setup, training) - Ongoing: $120K/year (subscription, support) - Time to value: 2-4 weeks to production deployment - 5-year TCO: $680K

91% TCO reduction over 5 years.

What It Delivers

For a mid-size manufacturer ($500M revenue, $50M inventory):

Operational improvements (conservative estimates): - **80-90% of routine planning decisions** handled without human intervention - **70% faster response** to supply chain disruptions (minutes vs. days) - **60% reduction in planner workload** on routine tasks (redirected to strategic work) - **25% inventory reduction** while maintaining or improving service levels - **20-35% total supply chain cost reduction** vs. current ordering policies

Financial impact: - $5M freed working capital (10% inventory reduction) - $1.875M/year holding cost savings - $5M revenue protection (20% stockout reduction) - $1.5M training cost savings (60% faster onboarding) - **$8.375M total annual operational value**

First-year ROI: 70x. Payback period: under one week.

Why Now

The pain points are reaching a breaking point across the industry:

- **Rising software costs**: Legacy vendor license renewals increasing 10-15% annually with no proportional improvement in capability
- **Implementation failures**: 40-60% of enterprise planning projects fail or exceed budget
- **Talent shortage**: Experienced planners retiring; the next generation expects modern tools, not green-screen ERP interfaces
- **Volatility as the norm**: Pandemics, trade wars, climate disruptions, and demand shifts have made deterministic planning demonstrably inadequate

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

- **AI pressure from above**: Boards and executives asking "why aren't we using AI?" while incumbent vendors offer limited, bolt-on capabilities at premium prices

Organizations that move first build a compounding advantage: their AI gets smarter, their teams get more productive, their institutional knowledge deepens — while competitors are still running weekly batch MRP and reviewing thousands of exceptions by hand.

Part 5: What Makes Autonomy Different

It's Not Another Planning Tool — It's an Operating Model Shift

Most planning software digitizes the existing process: take the spreadsheet, put it in a database, add a web interface. The planner still does the same work, just in a different window.

Autonomy changes who does what. The system does the routine planning. The human does the governance. This isn't a feature difference — it's a structural difference in how the planning function operates.

Trust Is Earned, Not Assumed

Legacy planning systems demand trust on day one. "Install our software, configure it for 12 months, and then trust the output." When planners don't trust it, they build shadow spreadsheets and override everything — which means the organization paid $6M for a system nobody uses.

Autonomy earns trust through evidence. Planners see AI recommendations alongside their own decisions, in identical scenarios, with measurable outcomes. They observe the reasoning. They test edge cases. They start with full control and gradually expand AI autonomy as confidence builds. By the time the system is running autonomously, trust is not an issue — it has been demonstrated hundreds of times.

The System Gets Smarter From Your Work

Every decision your team makes — and every override they provide — becomes training data. The system doesn't just execute; it learns. Not from generic industry data, but from your team's judgment, applied to your supply chain, in your operating context. This creates a flywheel:

```
 More decisions → Better AI → Less routine work →
More time for strategic decisions → Better strategic outcomes →
More trust → Expanded AI autonomy → More decisions handled → ...
```

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

After 12 months of operation, your system reflects the accumulated expertise of your entire planning team. This is a defensible competitive advantage that no amount of software licensing can replicate.

Risk-Free Testing of Any Change

The digital twin is not a gimmick — it's a decision-making tool. Before you change a safety stock policy, before you onboard a new supplier, before you close a warehouse, you run it through the twin. You see the impact across every node in the network, across hundreds of demand scenarios, with financial projections and confidence intervals. You make the decision with data, not intuition.

This alone justifies the platform for organizations that make even one significant supply chain structural change per year.

The sections that follow provide detailed technical architecture, AI model specifications, training workflows, integration capabilities, and implementation details for technical evaluation.

Revolutionary Continuous Autonomous Planning Architecture

Transforming Supply Chain Management from Cadence to Event-Driven

Traditional Planning (Weekly/Monthly Cadence):

```
 Monday:     Import SAP data
Tuesday:    Run MPS batch job
Wednesday:  Manual exception review
Thursday:   Approve and release plan
Friday:     Publish to ERP
Issues: 5-day latency, cannot react to intra-week disruptions
```

Continuous Autonomous Planning (Event-Driven with AI Agents):

```
 Event detected → Agent triggered → Replan affected SKUs → Commit to branch → Human notified
Benefits: Minutes from event to action, only affected products replanned, human-in-the-loop
```

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

The AIIO Framework: Automate-Inform-Inspect-Override

Core Philosophy: Agents autonomously handle routine planning within guardrails, inform humans of decisions, humans inspect results and override when necessary with captured context for continuous learning.

Stage	Responsibility	Who	Tool/Interface
Automate	Generate recommendations, execute within guardrails	AI Agents	Event-driven orchestration, A2A collaboration
Inform	Notify humans of decisions	AI Agents	LLM chat, email alerts, dashboard
Inspect	Review agent decisions, drill into details	Humans	LLM-first UI, point-and-click deep dive
Override	Change plans + provide learning context	Humans	UI with reason capture, file upload

Continuous Learning Loop:

```
 1. Automate  → Agents generate recommendations based on current models
2. Inform    → Notify humans of actions (auto-executed or pending approval)
3. Inspect   → Humans review via LLM chat or point-and-click UI
4. Override  → Humans provide feedback with context (reason, files, outcome preference)
5. Learn     → System captures context, identifies learning signals
6. Update    → Adjust guardrails, update KPI weights, fine-tune models
7. Improve   → Next iteration uses updated models/prompts/guardrails
→ Cycle repeats - agents get better over time
```

Learning Outcomes: - Week 1: 45% auto-execute rate, 35% override rate - Week 12: 72% auto-execute rate, 15% override rate - Target: 85% auto-execute rate, <10% override rate

Decision Stream — The Primary Interface

The **Decision Stream** (`/decision-stream`) is the default landing page for all users. Every supply chain decision — whether made by an AI agent or a human planner — appears here in real time as a decision card showing what was decided, the confidence level, and the status (Decided, Suggested, Overridden).

Pre-Computed Decision Reasoning: Every decision card includes an "Ask Why" button that expands inline to show the agent's reasoning instantly (<1ms). Reasoning is pre-computed at decision time from the actual signals the agent processed — not generated after the fact by an LLM. This covers all 11 TRM agents and all 3 GNN models

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

(S&OP GraphSAGE strategic risk, Network tGNN daily directives, Site tGNN hourly cross-TRM coordination).

Hive Context Capture: Each decision records the full signal bus state at the time of decision — urgency levels, triggering signals, cycle phase, and signals emitted — providing a complete audit trail of the multi-agent coordination context.

The traditional point-and-click planning UI (96+ pages covering demand planning, supply planning, inventory optimization, MPS, ATP worklists, etc.) remains available as a complementary deep-dive interface — but the Decision Stream is where planners spend most of their time.

AWS Supply Chain Integration

Insights → Event Triggers: - Stock-out risks, excess inventory alerts → Planning events - Vendor lead-time predictions → Supply disruption events - Forecast deviations → Demand change events

Recommendations → Agent Proposals: - Agents generate ranked recommendations using AWS SC metrics: - Risk resolution % (0-40 points) - Distance impact (0-20 points) - Sustainability (CO2 reduction, 0-15 points) - Service level impact (0-15 points) - Inventory cost (0-10 points)

Collaboration → A2A Communication: - Multi-agent negotiation (up to 3 rounds) - Conflict detection and resolution - LLM Supervisor mediation - Shared KPI optimization (service level, cost, inventory turns, CO2)

Complete AWS Supply Chain Feature Integration

Architectural Integration Strategy: - **Data Lake**: Databricks integration for ERP data ingestion and transformation - **Demand Planning**: External system integration with view-only visibility and delta analysis - **Supply Planning & Execution**: Full AWS SC-compliant planning and order management - **Insights & Risk**: ML-powered risk detection, watchlists, predictive analytics - **Recommendations**: AI-driven rebalancing engine with action scoring - **Collaboration**: Team messaging, commenting, and approval workflows

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

AWS SC Data Model Compliance (Backend - 100% Complete)

AWS SC Standard	Status	Implementation
Hierarchical Overrides	100%	6-level InvPolicy, 5-level VendorLeadTime, 3-level SourcingRules
Safety Stock Policies	100%	All 8 types: abs_level, doc_dem, doc_fcst, sl, sl_fitted, conformal, sl_conformal_fitted, econ_optimal
Vendor Management	100%	TradingPartner, VendorProduct, FK references
Sourcing Schedules	100%	Periodic ordering (weekly, monthly, custom)
Advanced Manufacturing	100%	Frozen horizon, setup/changeover, batch sizing, BOM alternates

See AWS_SC_100_PERCENT_COMPLETE.md for full certification details.

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

AWS SC Product Feature Parity (UI/UX - 100% Complete)

AWS SC Feature	Status	UI Coverage	Key Capabilities
Material Visibility	Complete	100%	Shipment tracking, ATP/CTP calculations, inventory projection, N-tier visibility with inventory flow, capacity analysis, risk assessment per tier
Order Planning & Tracking	Complete	100%	Unified all-order dashboard (PO/TO/CO/MO), order detail with timeline, vendor scorecard, KPI cards, search/filter, CSV export
Insights & Risk Analysis	Complete	100%	Risk detection, watchlists, alert rules configuration (create/edit/delete rules with metrics, thresholds, escalation), notification channels (email/in-app/Slack with enable/disable), predictive analytics
Recommended Actions	Complete	100%	Risk-based recommendations, scoring breakdown (5 components /100), accept/reject workflow, batch operations, simulation, performance tracking (acceptance rates, predicted vs actual outcomes, savings realization, type-level analytics)
Collaboration	Complete	100%	A2A/H2A/H2H framework, approval workflows, inline comments, activity feed, team messaging, notification preferences (per-alert-type, quiet hours, digest frequency, delivery channels)
Demand Planning	Complete	100%	Demand plan view/edit with wired import/export (CSV upload/download), forecast adjustment history, version comparison, consensus planning with voting, CPFR with exception detection, ML forecast pipeline
Data Lake	External	N/A	Using Databricks for ERP connectivity and data transformation

1. Insights & Risk Analysis

ML-Powered Risk Detection: - Automatic identification of stock-out and overstock risks - Probabilistic risk scoring (0-100) based on historical patterns - Early warning alerts (7-30 days advance notice) - Root cause analysis using decision trees and SHAP values

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Customizable Watchlists: - Critical product monitoring with custom thresholds - Multi-dimensional filters (product, location, supplier, customer) - Alert escalation rules (email, SMS, dashboard notification) - Watchlist templates for common scenarios (seasonality, new product launches, supplier risk)

Predictive Analytics: - Stock-out probability forecasts (7/14/30/90-day horizons) - Vendor lead-time prediction with confidence intervals - Demand surge detection (3-sigma anomaly detection) - Forecast vs. actual tracking with MAPE/RMSE metrics

Real-Time Alert System (Implemented): - Event-driven alerts triggered by threshold violations - Priority-based routing (P0: immediate, P1: 1-hour, P2: 4-hour, P3: daily) - Multi-channel delivery (dashboard, email, Slack, Teams) - Alert grouping and deduplication

Alert Rules Configuration (Implemented): - Create, edit, and delete custom alert rules with configurable metrics and thresholds - Per-rule severity assignment (CRITICAL/HIGH/MEDIUM/LOW) and alert type selection - Condition-based triggering (metric above/below threshold) - Escalation timer (auto-escalate unresolved alerts after configurable hours) - Per-rule notification channel selection (email, in-app, Slack) - Enable/disable individual rules without deletion

Notification Channels (Implemented): - Email, in-app, and Slack channel configuration with per-channel enable/disable - Webhook configuration for Slack integration - Channel status monitoring (active/disabled) - Per-alert-type notification preferences in Settings (stockout, overstock, lead time, recommendations, escalations, order exceptions) - Quiet hours with configurable start/end times (critical alerts bypass quiet hours) - Email digest frequency control (real-time, hourly, daily, weekly)

Uncertainty Quantification & Conformal Prediction (Implemented January 2026):

Distribution-free uncertainty quantification for supply chain planning with guaranteed prediction intervals:

Conformal Prediction: - Calibrate from historical Plan vs. Actual data (any variable: demand, lead time, yield) - Generate prediction intervals with formal coverage guarantees (no distribution assumptions) - Per-product, per-site calibration for precise uncertainty estimates - Adaptive conformal prediction with drift detection and automatic recalibration - Safety stock calculation with guaranteed service levels using conformal bounds

Key Benefits vs. Traditional Approaches: - Traditional: Assume normal distribution → intervals often wrong - Conformal: Data-driven intervals → guaranteed coverage (e.g., "90% coverage" means actual ≥90%)

Planning Method Comparison: - Side-by-side stochastic vs. deterministic analysis - Monte Carlo simulation with 1000+ scenarios - Sensitivity analysis for key parameters - Decision recommendations based on uncertainty level

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Capabilities: `view_risk_analysis`, `manage_watchlists`, `view_predictions`, `configure_alerts`, `view_uncertainty_quantification`

2. Recommended Actions Engine (AI Agent Suggested Actions)

AI-Powered Recommendations: The platform's three AI agents (LLM, GNN, TRM) continuously analyze supply chain state and generate suggested actions for human planners. Through multi-agent consensus with adaptive weight learning, the system provides high-confidence recommendations that improve over time.

Agent-Generated Rebalancing Recommendations: - Network-wide inventory rebalancing using linear programming (Optimizer Agent) - Multi-site transfer suggestions based on demand variance (GNN Agent with temporal pattern recognition) - Safety stock redistribution to minimize total network inventory (LLM Agent strategic planning) - Seasonal rebalancing for pre-positioning inventory (ML-Forecast Agent with time series analysis)

Multi-Agent Consensus Process: 1. **Individual Agent Analysis**: Each agent (LLM, GNN, TRM) independently analyzes the supply chain state 2. **Weighted Voting**: Agent recommendations combined using learned weights (e.g., LLM: 45%, GNN: 38%, TRM: 17%) 3. **Confidence Scoring**: Agreement between agents indicates recommendation reliability 4. **Adaptive Learning**: Weights automatically adjust based on historical performance

Action Scoring Algorithm: - **Risk Resolution** (0-40 points): % reduction in stock-out/overstock probability - **Distance Impact** (0-20 points): Transportation cost and lead time - **Sustainability** (0-15 points): CO2 reduction from optimized routing - **Service Level Impact** (0-15 points): Fill rate and OTIF improvement - **Inventory Cost** (0-10 points): Holding cost reduction

Impact Simulation & What-If Analysis: - Simulate recommended actions before execution - Compare multiple scenarios side-by-side (including AI vs. human decisions) - Monte Carlo simulation for uncertainty quantification - **Scenario-Based Testing**: Run suggested actions through digital twin simulations before production deployment - Rollback capability for risk-free testing

Decision Tracking & ML Learning Loop: - Track accepted/rejected recommendations - Capture human override reasons and context (RLHF - Reinforcement Learning from Human Feedback) - Fine-tune AI agents on human simulation and production data - Adaptive guardrails that learn from human decisions

Performance Tracking (Implemented): - Acceptance rate, override rate, and auto-execution rate KPIs - Performance breakdown by recommendation type (rebalance, expedite, inventory buffer) - Predicted vs actual savings comparison with realization rate percentage - Predicted vs actual service level impact comparison - Score effectiveness analysis (avg score of accepted vs overridden recommendations) - Recent outcomes

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

table with effectiveness assessment per recommendation - Total savings realized vs predicted with realization rate

Capabilities: `view_recommendations`, `manage_recommendations`, `approve_actions`, `simulate_impact`

3. Collaboration & Team Coordination

Team Messaging Interface (Implemented): - Channel-based messaging (group, direct, entity-linked to PO/TO/supply plans) - Threaded conversations with reply counts and thread view panel - @mentions with autocomplete and read tracking - Message pinning, editing, and deletion with audit trail - Read receipts per channel with unread count badges - Embedded in Collaboration Hub as dedicated tab

Inline Comments on Orders & Plans (Implemented January 2026): - Comment on purchase orders, transfer orders, supply plans, recommendations - Comment threading with nested replies - @mentions with user autocomplete and notifications - Comment types: general, question, issue, resolution, approval, rejection - Pin important comments for visibility - Edit/delete with audit trail - Real-time updates via WebSocket

Notification System: - Configurable notification preferences per user - Digest emails (real-time, hourly, daily) - Mobile push notifications - In-app notification center with priority filtering

Activity Feed & Audit Trail: - Chronological feed of all planning actions - User attribution and timestamp tracking - Filterable by action type, user, date range - Export to CSV for compliance reporting

Document Sharing: - Attach files to orders, plans, comments - Supported formats: PDF, Excel, Word, images - Version control for document updates - Permission-based access control

Capabilities: `view_collaboration`, `post_messages`, `manage_collaboration`, `view_activity_feed`

4. Order Planning & Tracking

Comprehensive Order Types: - **Purchase Orders (PO)**: Vendor procurement with approval workflows - **Transfer Orders (TO)**: Inter-site transfers with transportation tracking - **Manufacturing Orders (MO)**: Production scheduling with BOM explosion - **Project Orders**: Custom project-based fulfillment - **Maintenance Orders**: Repair and maintenance scheduling - **Turnaround Orders**: Refurbishment and reverse logistics

Order Lifecycle Management: - Draft → Submitted → Approved → Sent → Acknowledged → Confirmed → Shipped → Received → Closed - Automatic state transitions based on business rules - Manual override capability for exception handling - Configurable approval workflows (single-level, multi-level, matrix)

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

PO Acknowledgment & Goods Receipt (Implemented January 2026): - Send PO to supplier with tracking - Record supplier acknowledgment with expected delivery date - Record supplier confirmation (ready for shipment) - Goods receipt with partial receive support - Quality inspection: accept/reject with reason codes - Variance tracking: over/under delivery detection - Multiple receipts per PO (partial deliveries) - Automatic PO status updates (PARTIAL_RECEIVED, RECEIVED)

Multi-Site Order Coordination: - Cross-site visibility of order status - Dependent order tracking (parent-child relationships) - Consolidated order views by product, supplier, customer - Network-wide order promising with ATP/CTP

Unified All-Order Dashboard (Implemented): - Single view combining all order types (PO, TO, CO, MO) with type tabs and filtering - 8 KPI summary cards (total, in-transit, processing, delayed, delivered, total value, avg priority, on-time %) - Full-text search by order ID or product, status filter dropdown - Order detail modal with timeline visualization (created → confirmed → in-transit → delivered) - CSV export of filtered order data - Vendor performance scorecard (on-time delivery %, quality rate, avg lead time, trend indicator)

Capabilities: `view_order_planning`, `create_order`, `approve_order`, `manage_project_orders`

5. Material Visibility & Shipment Tracking

Real-Time Shipment Tracking: - Integration with logistics providers (FedEx, UPS, DHL, 3PL) - GPS tracking for in-transit inventory - Estimated delivery time with delay alerts - Exception management for at-risk shipments

Inventory Projection & ATP/CTP: - Available-to-Promise (ATP): On-hand + scheduled receipts - Capable-to-Promise (CTP): ATP + planned production - Multi-level ATP/CTP with BOM explosion - Time-phased projections (daily/weekly buckets)

Delivery Risk Analytics: - Probabilistic on-time delivery prediction (0-100%) - Risk factors: weather, port congestion, carrier reliability - Mitigation recommendations: expedite, reroute, safety stock - Proactive customer notification for at-risk orders

N-Tier Network Visibility (Complete): - Multi-echelon supply chain visualization (supplier → DC → customer) - DAG-based network topology with Sankey diagrams - Pipeline inventory tracking (in-transit, in-production) - Bottleneck identification and capacity constraints - **Inventory Flow**: Tier-to-tier flow visualization with volume, in-transit counts, transit time, lane status (normal/constrained), inventory levels vs targets comparison table - **Capacity Analysis**: Per-tier utilization bars with color-coded thresholds (green <80%, amber 80-90%, red 90%+), bottleneck identification, detailed metrics table (throughput, quality, OTD), optimization recommendations - **Risk Assessment**: Per-tier risk scores (0-100) with donut visualization, 5-category risk

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

breakdown (supply continuity, capacity, quality, lead time, financial), mitigation action table with priority and due dates, all-mitigations summary sorted by priority

Capabilities: `view_shipment_tracking`, `manage_shipments`, `view_inventory_visibility`, `view_ntier_visibility`

6. Demand Planning (Full-Featured)

Demand Plan View & Edit: - Display current demand plan by product/location/time with P10/P50/P90 confidence intervals - Inline cell editing with undo/redo support - Bulk adjustments (percentage, delta) across multiple forecasts - Filter by product family, location hierarchy, date range - Import/Export for bulk operations (Wired: CSV upload via file picker, CSV download with template fallback, import validation with error feedback)

Forecast Adjustment History (Implemented March 2026): - Full audit trail of all forecast adjustments with filtering - Filter by reason code (promotion, seasonal, event, market intelligence, correction) - Filter by source (manual, bulk, agent, import) and status (applied, pending, reverted) - One-click revert capability for applied adjustments - Summary statistics: total, applied, reverted, pending approval

Version Comparison (Implemented March 2026): - Side-by-side comparison of any two forecast versions - Delta summary: products changed, average change %, increases vs decreases - Product-level, site-level, and period-level difference table - Create snapshots (snapshot, baseline, consensus, published types) - Locked version support for approved baselines

Consensus Planning (Implemented March 2026): - Multi-stakeholder forecast consensus cycles (DRAFT → COLLECTING → REVIEW → VOTING → APPROVED → PUBLISHED) - Source-based submissions: sales, marketing, finance, operations, statistical - Phase-based progress tracking with visual pipeline - Voting mechanism: approve, reject, abstain, request changes - Version comparison across stakeholder submissions - Threaded discussion per consensus cycle

ML Forecast Pipeline: - Configurable ML pipeline with 30+ parameters - Clustering-based statistical forecasting - Pipeline run management with async execution - Publish results to production forecasts

CPFR (Demand Collaboration) (Implemented March 2026): - Trading partner collaboration with forecast sharing - Exception detection (>20% variance threshold) - Approval/rejection workflows with reason capture - Forecast accuracy tracking per partner (MAPE, bias, hit rate) - Collaboration types: forecast share, consensus, alert, exception

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Forecast Exceptions: - Exception detection and management with configurable rules - Variance threshold, trend detection, outlier detection, bias detection - Severity levels and resolution workflows - Comment threading per exception

Capabilities: `view_demand_planning`, `manage_demand_planning`, `view_demand_collaboration`, `view_forecasting`, `view_forecast_exceptions`

Key Innovations of Continuous Autonomous Planning

1. Event-Driven Agent Orchestration: - Agents react to real-time events (orders, disruptions, forecast changes) - Priority-based task queue (P0 < 1 min, P1 < 5 min, P2 < 1 hour, P3 < 24 hours) - Agent dependency management (Policy agents → Execution agents → Supervisor)

2. Git-Like Plan Versioning (Kinaxis-Inspired):

```
main (production plan)
├── daily/2026-01-23  (today's baseline)
│   ├── scenario/mps-order-789  (MPS agent scenario)
│   ├── scenario/supplier-delay-456  (MRP agent scenario)
│   └── hotfix/stockout-case  (Emergency fix)
```

- Branch/merge workflow for plan changes
- Full plan history with commit/diff/revert
- Scenario-based "what-if" analysis

3. Incremental CDC Snapshotting: - Nightly full snapshot: 260 MB (50 MB compressed) - Hourly incremental: 10-50 KB (100x reduction) - Agent-triggered incremental: Real-time CDC - 28-day retention: ~83 GB total vs. 26 GB if full snapshots hourly

4. LLM-First UI:

```
 Planner: "Show me all agent actions from today"
LLM: "I found 12 agent actions today:
     - 5 by MPS Agent (all auto-executed, avg cost impact: +$1,200)
     - 3 by MRP Agent (2 auto-executed, 1 awaiting your review)
     [View All] [Filter by Impact] [Show Pending Only]"

Planner: "Show me the MRP action that needs my review"
LLM: "The MRP Agent recommends creating a PO for 500 units of BOTTLE...
     This requires review because order value ($62,000) exceeds threshold ($50,000)
     [Approve] [Modify] [Request Alternatives]"
```

5. Guardrails for Agent Autonomy: - Financial: Max cost increase 5%, max PO value $50K - Operational: Max safety stock increase 20%, max production change 15% -

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Customer: Min service level 92%, max order delay 3 days - Strategic: Max CO2 increase 10%, min supplier diversity 2

6. Human Override with Context Capture: When humans override agents, they provide: - **Reason category** (10 options): Domain knowledge, external constraint, strategic decision, data quality, policy exception, customer priority, supply chain insight, cost sensitivity, risk aversion, other - **Free-text explanation** (min 10 characters) - **Supporting files**: PDFs, emails, spreadsheets → System learns: Adjust guardrails, update KPI weights, fine-tune models, update LLM prompts

7. Continuous Order Promising: - **Batch Promising**: Priority-aware allocation (VIP customers first) - **Continuous Promising with Batched Notification** (Recommended): - Promise on arrival, batch notifications - High-priority orders can pre-empt lower-priority promises before notification - Priority override window: 30 minutes configurable

Powell Sequential Decision Analytics Framework

The platform's planning architecture is grounded in Warren B. Powell's **Sequential Decision Analytics and Modeling (SDAM)** framework, which provides a unified theoretical foundation for decision-making under uncertainty.

Powell's Five Core Elements:

Element	Symbol	Platform Implementation
State	S_t	Inventory levels, backlog, pipeline, demand history, network topology
Decision	x_t	Order quantities, production schedules, sourcing allocations
Exogenous Information	W_{t+1}	Customer demand, lead time realizations, yield variations
Transition Function	S^M	Supply chain simulation engine (BeerLine, SimPy)
Objective Function	F	Minimize E[total cost] subject to service level constraints

Powell's Four Policy Classes:

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Class	Acronym	Description	Platform Use
Policy Function Approximation	PFA	Direct S→x mapping	Base-stock rules, historical proportions
Cost Function Approximation	CFA	Parameterized optimization	S&OP/MPS with policy parameters θ
Value Function Approximation	VFA	Q-learning, TD learning	TRM agent, Execution tGNN
Direct Lookahead	DLA	Model predictive control	Strategic planning with scenarios

Hierarchical Policy Integration:

The platform implements Powell's insight that **multi-level planning is nested optimization**:

```
 Strategic (DLA): θ_strategic = argmax E[V(S) | scenarios]
                        ↓ bounds
S&OP (CFA):     θ_sop = argmin Cost(θ) s.t. θ ∈ θ_strategic
                        ↓ parameters
MPS (CFA):      θ_mps = optimize(production) s.t. θ ∈ θ_sop
                        ↓ constraints
Execution (VFA): x* = argmax Q(s,a|θ_mps)
```

Key Powell Insights Implemented:

1. **Aggregation as State Abstraction**: Higher planning levels (S&OP) work with aggregated state representations to reduce dimensionality. The `AggregationService` rolls up Site→Country, SKU→Family, Day→Month.

2. **Disaggregation as Policy**: Powell treats disaggregation proportions as **policy decisions**, not fixed transformations:

3. **PFA**: Fixed historical proportions (simple, stable)
4. **CFA**: Learned proportions optimized from data (adaptive)

5. **VFA**: Value-based allocation maximizing downstream value

6. **Hierarchical Consistency**: Lower-level value functions must approximate upper-level expectations: `V_execution ≈ E[V_tactical | disaggregate(S_tactical)]` Platform enforces <10% deviation tolerance between planning levels.

7. **Monte Carlo for Evaluation, Optimization over Scenarios**: Current stochastic planning generates scenarios and evaluates outcomes. Powell recommends **extracting optimal policy parameters** from scenario results, which the platform implements through adaptive weight learning.

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Reference: See POWELL_APPROACH.md for complete implementation details.

Agent Hierarchy and Tasks

Policy Agents (Set the rules): - Inventory Policy Agent: Safety stock, reorder points, DOC targets - Sourcing Policy Agent: Make-vs-buy, vendor selection, priorities - Capacity Policy Agent: Shift patterns, overtime rules, bottlenecks - Demand Policy Agent: Forecast models, seasonality, demand sensing

Execution Agents (Implement the plan): - MPS Agent: Master production schedule - MRP Agent: Component requirements, purchase requisitions - Capacity Agent: Resource scheduling, load balancing - Order Promising Agent: ATP/CTP calculation, order confirmation

Supervisor Agents (Coordinate and escalate): - LLM Supervisor: Exception handling, root cause analysis, human escalation - Global Planner Agent: Network-wide optimization, trade-off analysis

Agent Collaboration Example:

```
 Policy Change Event: "Increase CASE safety stock from 50 to 75"
    ↓
Inventory Policy Agent recalculates targets for all CASE-related products
    ↓
MPS Agent replans affected weeks to meet new targets
    ↓
MRP Agent adjusts component requirements based on new MPS
    ↓
Capacity Agent verifies resource availability for new MRP plan
    ↓
LLM Supervisor reviews cascade impact, alerts planners if significant
```

Performance & Scalability

Event Processing Capacity: - 100,000 events/day (peak: 10 events/sec) - 50,000 agent tasks/day - 200 incremental commits/day - 1,000 plan diff queries/day

Latency Targets: - P0 (Critical): < 1 minute event → plan → publish - P1 (High): < 5 minutes - P2 (Medium): < 1 hour - P3 (Low): < 24 hours

Storage Requirements (1 year): - Full snapshots: 18 GB - Incremental snapshots: 3.6 GB - Events log: 36 GB - Agent decisions: 18 GB - **Total: ~83 GB/year**

Cost Analysis (AWS): - Infrastructure: ~$632/month - Cost per agent decision: $0.0094 (or $0.082 with LLM review)

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Continuous Autonomous Planning vs. Legacy Systems

Aspect	Legacy (Kinaxis/SAP IBP)	Continuous Autonomous Planning Platform
Planning Frequency	Weekly/monthly batch	Event-driven (minutes)
Latency	5-7 days event → action	<5 minutes (P0/P1 events)
Human Involvement	Manual batch review of 1000s SKUs	Inspect exceptions only (~10-20/day)
Plan Changes	Full replan (slow)	Incremental (only affected SKUs)
Versioning	Single "current" plan	Git-like branches, full history
AI Integration	Limited, bolt-on	Native agents with A2A collaboration
Learning	Consultant retraining	Continuous from human overrides
Explainability	Black box optimization	Pre-computed reasoning on every decision (<1ms Ask Why)
UI Paradigm	Complex ERP screens	Decision Stream (primary) + point-and-click deep dive
Cost	$100K-$500K/user/year	$10K/user/year (90% reduction)
Deployment	12-18 months	2-4 weeks

Value Proposition: Three Systems in One

1. Enterprise Planning Platform (90% cost reduction vs. Kinaxis/SAP IBP) - Continuous event-driven planning with AI agents - Git-like versioning and scenario management - LLM-first UI with natural language queries - Incremental CDC snapshotting for efficiency - Plan vs. actual automated comparison

2. AI Agent Training Ground (Risk-free validation before production) - Test agent strategies in simulation environments - Build stakeholder confidence through competitive simulation - Generate diverse training data from human decisions - RLHF (Reinforcement Learning from Human Feedback)

3. Stochastic Simulation Engine (Model real-world uncertainty) - 20 distribution types for operational variability - Monte Carlo simulation (1000+ scenarios) - Probabilistic balanced scorecard (P10/P50/P90 KPIs) - Multi-echelon DAG topology support

4. Conformal Prediction Engine (Implemented January 2026) - **Distribution-Free Intervals**: Guaranteed prediction coverage without assuming normal/lognormal distributions - **Calibration from History**: Use Plan vs. Actual data to calibrate per-product, per-site uncertainty bounds - **Formal Guarantees**: If we promise "90%

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

coverage", actual coverage will be ≥90% (mathematically proven) - **Safety Stock with Guarantees**: Calculate safety stock with formal service level guarantees - **Adaptive Recalibration**: Automatic drift detection and recalibration when forecast accuracy degrades - **Planning Method Comparison**: Side-by-side stochastic vs. deterministic with decision recommendations

5. Distribution Fitting & Likelihood Estimation Engine (Implemented February 2026) - **MLE Distribution Fitting**: Automatically fit 20 distribution types (Normal, Lognormal, Gamma, Weibull, Beta, Exponential, Triangular, Mixture, etc.) to historical demand, lead time, yield, and price data using maximum likelihood estimation - **Statistical Model Selection**: Kolmogorov-Smirnov goodness-of-fit testing, AIC/BIC ranking to select the best-fitting distribution per product-site combination - **Distribution-Aware Safety Stock** (`sl_fitted` policy): When demand or lead time is non-Normal (e.g., lognormal, Weibull), uses Monte Carlo simulation of Demand-During-Lead-Time (DDLT) instead of the standard z-score formula — correcting 15-30% safety stock miscalculation common with skewed distributions - **Distribution-Aware Feature Engineering**: TRM state vectors include fitted distribution parameters (shape, scale, skewness, kurtosis) and MAD/median ratios instead of mean/std, following Kravanja (2026) on robust feature design for non-Normal data - **API**: `POST /api/v1/stochastic/fit` for on-demand distribution fitting with ranked results - **Reference**: See TECHNICAL_OVERVIEW.md Part 8 for full architecture

6. Digital Twin Training Pipeline (Implemented February 2026) - **Six-Phase Cold-Start**: Takes AI agents from zero experience to production autonomy in 3-5 weeks: 1. **Behavioral Cloning** (hours): Each TRM matches the deterministic engine baseline within ±5% 2. **Coordinated Simulation** (2-3 days): All 11 TRMs train simultaneously with signal bus active, learning inter-agent coordination from 28.6M+ training records 3. **Site tGNN Training** (~1 day): Cross-TRM coordination model learns causal relationships from coordinated traces, then fine-tunes via RL with site-level balanced scorecard as reward 4. **Stochastic Stress-Testing** (3-5 days): TRMs + Site tGNN together face adversarial scenarios — 3σ demand spikes, supplier failures, capacity shocks, compound disruptions 5. **Copilot Calibration** (2-4 weeks): Agents run in copilot mode with human overrides captured and weighted by Bayesian effectiveness tracking; Site tGNN in shadow mode 6. **Autonomous CDC Relearning** (ongoing): Continuous improvement from production outcomes via the CDC → Relearning loop; Site tGNN retrained every 12h - **Warm-Start Advantage**: Production AI starts with pre-trained knowledge from simulation, not random initialization — eliminating the months-long "cold start" period that makes enterprise AI adoption impractical - **Reference**: See TRM_HIVE_ARCHITECTURE.md Section 15 for implementation details

7. Vertical Escalation (Escalation Arbiter) (Implemented March 2026) - **Problem**: Some execution-level anomalies signal that strategic policy parameters are wrong, not that execution decisions need fine-tuning. Retraining TRMs doesn't fix a safety stock multiplier that's too low. - **Persistence Detection**: Monitors TRM decision patterns

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

across all sites for persistent directional drift — when agents consistently correct in the same direction (e.g., always ordering 20% more), the policy is wrong - **Intelligent Routing**: Routes escalations to the right level — horizontal (CDC retrain), operational (off-cadence tGNN refresh), or strategic (S&OP policy review) — based on pattern scope (single TRM vs. multi-site vs. network-wide) - **Theoretical Foundation**: Maps to Kahneman's dual-process theory (TRMs = System 1 fast thinking, tGNN/GraphSAGE = System 2 slow thinking) and Boyd's nested OODA loops (execution/operational/strategic time horizons) - **Reference**: See ESCALATION_ARCHITECTURE.md for full theoretical framework

8. Quantitative Supply Chain Economics (Implemented March 2026) - **Dollar-Denominated Decisions**: AI agent reward functions now use actual economic costs (holding cost per unit per day, stockout cost per unit, ordering cost per order) computed from product unit cost and inventory policy parameters — replacing heuristic scaling factors - **Economically Optimal Safety Stock** (`econ_optimal` policy): Marginal economic return analysis — stock one more unit only when the expected stockout cost exceeds the holding cost. Monte Carlo simulation over fitted demand and lead time distributions with no fallbacks or defaults - **CRPS Forecast Scoring**: Continuous Ranked Probability Score — the gold standard for probabilistic forecast evaluation — integrated into the conformal prediction pipeline. Tracks forecast distribution quality, not just coverage - **Censored Demand Detection**: Automatic detection of stockout periods where observed demand understates true demand. Censored observations excluded from distribution fitting to prevent systematic underestimation - **Log-Logistic Lead Times**: Fat-tailed log-logistic distribution added for lead time fitting, capturing the "most shipments arrive on time, but some arrive very late" pattern better than Weibull or lognormal - **Automated Policy Re-Optimization**: Weekly CFA (Cost Function Approximation) job using Differential Evolution to re-optimize inventory policy parameters (θ), ensuring policies stay current as demand patterns evolve - **No Fallbacks**: All economic parameters are explicitly required for every tenant — the system raises errors rather than silently using defaults. Every cost assumption is visible and auditable - **Reference**: See TECHNICAL_OVERVIEW.md Part 11 for architecture details

10. Learned Intra-Site Cross-TRM Coordination (Site tGNN) (Implemented March 2026) - **The Gap**: Between sub-millisecond reactive stigmergic signals (Layer 1) and daily network-wide inference (Layer 3), many cross-TRM interactions follow predictable causal chains that neither mechanism captures. A production spike today will generate quality inspection pressure in hours and maintenance load tomorrow -- but the signal bus only reacts after each link materializes. - **Layer 1.5**: A lightweight graph neural network (~25K parameters, GATv2+GRU) that learns causal relationships between the 11 execution agents within each site, modeling 22 directed causal edges (ATP-to-MO, MO-to-Quality, Quality-to-Buffer, etc.). Predicts cross-TRM cascade effects that reactive signals alone cannot capture. - **Urgency Modulation**: Outputs small adjustment deltas ([-0.3, +0.3]) that modulate the UrgencyVector before the 6-phase decision cycle runs.

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

The Site tGNN shifts emphasis across the TRM hive -- "pay more attention to quality today because production volume is spiking" -- rather than overriding individual TRM decisions. - **Training**: 3-phase pipeline (behavioral cloning from MultiHeadTrace records, PPO fine-tuning with site-level Balanced Scorecard reward, production calibration in shadow mode). Feature-flagged OFF by default; zero adjustments when untrained. - **Reference**: See TECHNICAL_OVERVIEW.md Part 13 and POWELL_APPROACH.md Section 5.20

9. Decision Intelligence Platform Architecture (Implemented, Gartner-aligned March 2026) - **Decision-Centric Architecture**: Gartner designated Decision Intelligence as "transformational" (2025 AI Hype Cycle) and published its inaugural Magic Quadrant for Decision Intelligence Platforms in January 2026. The platform implements the full DI lifecycle: decision modeling (Powell SDAM), orchestration (TRM Hive + AAP), monitoring (CDC + conformal prediction + CRPS), and governance (override tracking + CDT risk bounds + escalation arbiter) - **Decisions as Assets**: Every recurring supply chain decision — stocking, ordering, rebalancing, allocation — is modeled with defined inputs, logic, constraints, ownership, and measured outcomes across 11 specialized agent types - **Three-Level Maturity Progression**: Decision Support (human in the loop) → Decision Augmentation (human on the loop, copilot mode) → Decision Automation (human out of the loop, autonomous mode) — with progression governed by measured decision quality, not arbitrary trust thresholds - **Gartner Four Use Cases**: Decision Stewardship (override governance + audit), Decision Analysis (8 policy types + probabilistic BSC), Decision Engineering (11 TRM agents + AAP orchestration), Decision Science (hybrid TRM + Claude Skills + conformal routing) - **Agentic AI Alignment**: Gartner predicts 50% of SCM solutions will use intelligent agents by 2030. Autonomy already deploys 11 specialized agents per site with multi-site coordination — positioning years ahead of the market curve - **Reference**: See TECHNICAL_OVERVIEW.md Part 12 and Decision Intelligence Framework Guide

Business Impact: Continuous Autonomous Planning ROI

Operational Benefits: - **80-90% automation rate** for routine planning tasks (agents handle without human review) - **70% faster response** to supply chain disruptions (minutes vs. days) - **60% reduction** in planner workload (inspect exceptions only, not full batch review) - **40% improvement** in forecast accuracy (continuous learning from overrides) - **25% inventory reduction** while maintaining service levels (optimized safety stock) - **20-35% cost reduction** vs. naive policies (AI agent optimization)

Financial Impact (Mid-size manufacturer, $500M revenue): - **Direct Cost Savings**: $5.8M first year vs. Kinaxis/SAP IBP - **Operational Benefits**: $8.375M/year (inventory reduction + holding cost savings + stockout reduction) - **Total First-Year Value**: $14.175M - **Platform Cost**: $200K - **ROI**: 70.9x return, 5.3-day payback period

Strategic Advantages: - **No vendor lock-in**: Open architecture, data portability - **Fast deployment**: Weeks instead of 12-18 months - **Transparent AI**: Build trust before

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

production deployment - **Continuous improvement**: Agents learn from human feedback - **Flexibility**: DAG-based topology, any network structure

Platform Capabilities

1. Continuous Autonomous Planning: Event-Driven Supply Chain Optimization

Event-Driven Planning Flow: 1. **Event Detection**: AWS SC insights, SAP CDC, manual triggers 2. **Agent Orchestration**: Priority queue routes events to agents 3. **A2A Collaboration**: Multi-agent negotiation for shared KPIs 4. **Scenario Generation**: Agents create plan branches with recommendations 5. **Guardrail Check**: Auto-execute if within bounds, else human approval 6. **Human Notification**: LLM chat + email + dashboard alerts 7. **Plan Commit**: Git-like commit to scenario branch 8. **Human Inspection**: Natural language queries + point-and-click deep dive 9. **Human Override** (if needed): Provide context (reason, files, outcome preference) 10. **Learning Pipeline**: Adjust guardrails, update KPIs, fine-tune models 11. **Merge to Main**: Execute approved plan 12. **Publish to ERP**: Incremental changes flow to SAP/Oracle

Agent-Generated Recommendations: - Ranked by score (0-100): Risk resolution + distance + sustainability + service level + cost - Simulated in scenario branches (what-if analysis) - Impact estimation: Before/after KPIs with confidence intervals - Alternatives considered: Top 3 options with trade-off analysis - Natural language explanation: "Why did I recommend this?"

Human Override Context Capture:

```
 Override Reason (Required)
Category: [External Constraint ✓]
Details: Vendor-B is undergoing facility audit, cannot accept orders for 2 weeks
Supporting Files:   vendor_b_audit_notice.pdf (125 KB)
Alternative Action: Use Vendor-C instead (14-day lead time vs. 10-day)
Impact: Service level 96.8% → 95.5% (-1.3%), Cost +$1,520

→ System learns: Add constraint "Vendor-B unavailable 2026-01-23 to 2026-02-06"
                 Adjust scoring to prefer Vendor-C during this period
                 Fine-tune LLM prompt with new external factor
```

2. Simulate: Interactive Supply Chain Simulation

Multi-Participant Scenarios: - 2-8 participants in roles (Retailer, Wholesaler, Distributor, Factory, Suppliers) - Real-time WebSocket updates - Period-by-period decision making - Instant cost and service level feedback

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Mixed Human-AI Scenarios: - Humans compete alongside or against AI agents - Multiple AI strategies: Naive, Conservative, Bullwhip, ML-Forecast, Optimizer, Reactive, LLM-powered - AI opponents provide consistent benchmarks for human performance - Learn from AI decisions in real-time

Scenario Variants: - **Default TBG**: Classic 4-echelon beer supply chain - **Three FG TBG**: Multi-product supply chain with shared components - **Variable TBG**: Dynamic demand patterns with high variability - **Complex SC**: Real-world DAG topologies with manufacturing BOMs - **Custom Configurations**: Create any supply chain structure via DAG editor

2. Plan: AI-Powered Supply Chain Optimization

Temporal Graph Neural Network (GNN): - Learns supply chain dynamics from historical data - Predicts demand propagation and bullwhip effects - Optimizes order quantities across entire network - Captures temporal dependencies (lead times, delays, seasonality) - Generalizes across different supply chain topologies

LLM-Based Multi-Agent System: - OpenAI GPT-powered agents for each supply chain role - Supervisor agent validates and improves recommendations - Global planner agent coordinates network-wide strategy - Natural language explanations of AI decisions - Adapts to changing conditions with zero retraining

Optimization Strategies: - Cost minimization (inventory holding + backlog penalties) - Service level maximization (fill rate targets) - Bullwhip effect reduction (demand variability dampening) - Multi-objective optimization (Pareto frontier exploration)

Planning Scenarios: - **What-If Analysis**: Simulate demand shocks, supplier disruptions, lead time changes - **Capacity Planning**: Test network under growth scenarios - **Risk Assessment**: Identify vulnerable sites and failure modes - **Strategic Design**: Optimize network topology (where to hold inventory, safety stock levels)

3. Analyze: Performance Intelligence & Benchmarking

Current Performance Analysis: - Upload historical scenario data or operational metrics - Compare actual decisions against AI-optimal strategies - Identify cost reduction opportunities - Quantify bullwhip effect in real operations - Benchmark against industry best practices

Advanced Analytics Dashboard: - **Bullwhip Metrics**: Demand amplification ratios across echelons - **Cost Breakdown**: Holding costs vs. backlog costs by role - **Service Levels**: Fill rates, stockout frequencies, customer satisfaction - **Efficiency Scores**: Order variability, inventory turnover, cash-to-cash cycle - **Pattern Detection**: Seasonality, trends, anomalies in demand/supply

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Uncertainty Quantification (Implemented January 2026): - **Conformal Prediction**: Distribution-free prediction intervals with formal coverage guarantees - Calibrate on historical Plan vs. Actual data - Generate valid prediction intervals without distributional assumptions - P(actual ∈ interval) ≥ 1-α guaranteed (typically 90-95% coverage) - Adaptive conformal prediction for time series with distribution shift - Safety stock calculation with formal service level guarantees - **Planning Method Comparison**: Side-by-side Stochastic vs. Deterministic analysis - Monte Carlo simulation (stochastic): Full uncertainty propagation, risk quantification - Deterministic + Sensitivity Analysis: Single-point optimization with parameter variation - Automated recommendation for which approach suits your use case - Cost distribution comparison (P10/P50/P90/P95) - Service level probability analysis

Reporting & Insights: - Automated scenario reports with AI-generated recommendations - Export to CSV, JSON, Excel for external analysis - Integration-ready APIs for BI tools (Tableau, Power BI) - Executive summaries with KPIs and action items

Comparative Benchmarking: - Human vs. AI performance head-to-head - Team vs. team competition - Historical trend analysis (improving or degrading?) - Industry peer comparisons (anonymized leaderboards)

4. Simulation: Confidence Building & Agent Improvement

Building Confidence in AI Agents

The Trust Problem: Organizations struggle to adopt AI recommendations because: - "Black box" algorithms feel risky - No way to validate AI logic before deployment - Disconnect between data scientists and supply chain practitioners

The Digital Twin Solution: Scenarios as Fast-Forward Supply Chain Simulations:

Critical Insight: A "scenario" is not entertainment—it's a **digital twin of the real supply chain** that executes in fast-forward time with synthetic demand. The only differences between a scenario and production execution are: - **Time Scale**: Scenarios advance in seconds/minutes; production advances in days/weeks - **Demand Source**: Scenarios use synthetic demand patterns; production uses actual customer orders - **Everything Else is Identical**: Same planning logic, same AI agents, same decision-making, same cost calculations

Scenarios Enable Three Critical Functions:

1. **Adoption Through Acceptance** (Build Trust Before Production):
2. **Transparent Competition**: Humans compete against AI in identical scenarios
3. **Observable Decisions**: Watch AI order quantities period-by-period, understand its logic
4. **Measurable Outcomes**: AI wins? Loses? By how much? Why?

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

5. **Safe Learning Environment**: No real inventory at risk, infinite practice periods
6. **Gradual Adoption**: Start with AI suggestions (copilot mode), progress to AI autonomy as confidence grows
7. **Stakeholder Buy-In**: Executives see AI performance before approving production deployment
8. **Policy Testing** (Risk-Free What-If Analysis):
9. **Inventory Policy Changes**: Test new safety stock levels, reorder points, service level targets
10. **Ordering Policy Changes**: Test different replenishment strategies (base-stock, (s,S), periodic review)
11. **Multi-Agent Weight Tuning**: Test different agent consensus weights (e.g., LLM: 50% vs. GNN: 40%)
12. **Guardrail Calibration**: Test min/max order constraints, rush order thresholds
13. **Lead Time Strategies**: Test impact of expedited shipping, supplier changes
14. **Cost Trade-offs**: Test holding cost vs. shortage cost sensitivity
15. **Structural Testing** (Network Redesign Validation):
16. **Add/Remove Sites**: Test impact of new distribution centers, closing warehouses
17. **Supplier Changes**: Test multi-sourcing, backup suppliers, supplier reliability
18. **Capacity Changes**: Test production capacity increases, storage expansions
19. **Network Topology**: Test hub-and-spoke vs. direct-ship models
20. **BOM Changes**: Test make-vs-buy decisions, component substitutions
21. **Lead Time Changes**: Test nearshoring, offshore alternatives

Transfer Learning: Train in Scenarios, Deploy to Production: - Run 50-100 scenarios with different demand patterns → Learn optimal agent weights - Validate statistical significance ($p < 0.05$) → Build confidence in approach - Deploy learned weights to production → Continue adapting to real data - Result: Production AI starts with pre-trained knowledge, not random initialization

Confidence Metrics: - Win rate: AI vs. Human in fair competition (Target: >70% AI win rate) - Cost differential: How much does AI save vs. best human participant? (Target: 20-35% reduction) - Consistency: AI variance vs. human variance across scenarios (Target: <50% human variance) - Explainability: Can humans articulate why AI made specific decisions? (Measured via post-scenario surveys) - Acceptance Rate: % of AI suggestions accepted by humans in copilot mode (Target: >80% acceptance)

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Agent Improvement Through Simulation

Continuous Learning Pipeline:

1. **Humans Run Scenarios** → Generate diverse decision data
2. **AI Observes Human Strategies** → Learn from expert participants
3. **AI Competes Against Humans** → Validate and test improvements
4. **Humans Learn from AI** → Adopt successful AI patterns
5. **Cycle Repeats** → Continuous improvement for both

Reinforcement Learning from Human Feedback (RLHF): - Capture expert human decisions as training labels - Fine-tune AI agents on human decision data - Reward agents for outperforming human benchmarks - Penalize agents for suboptimal or unstable behavior

Achievement-Driven Training: - 17 achievements across 5 categories (Cost Control, Service Excellence, Collaboration, Innovation, Mastery) - Each achievement represents a desirable supply chain behavior - Train AI agents to unlock achievements (simulation reward functions) - Achievement completion rate becomes agent performance metric

Example Achievement → Agent Objective Mapping: - **"Steady Hand"** (keep inventory variance low) → Train agent to minimize order variability - **"Bullwhip Tamer"** (reduce amplification) → Reward stable ordering policies - **"Zero Backlog"** (no stockouts) → Optimize service level targets - **"Negotiation Master"** (collaboration) → Multi-agent coordination rewards

Simulation ROI: - **Engagement**: 3-5x higher participation in training vs. traditional methods - **Retention**: 70% knowledge retention after 6 months (vs. 20% for lectures) - **Performance**: Participants improve 40-60% faster with simulation feedback - **Agent Quality**: Human-validated agents have 25% lower deployment failure rate

---

Technical Architecture

Multi-Echelon DAG Framework

DAG-Based Supply Chain Topology: - Directed Acyclic Graph (DAG) represents material flow - 4 master site types: Market Supply, Market Demand, Inventory, Manufacturer - Flexible configuration: serial chains, convergent networks, divergent distribution - Supports Bill of Materials (BOM) transformations (e.g., 1 Case = 4 Six-Packs)

Variability Modeling: - **Demand Variability**: Stochastic demand with configurable distributions (normal, uniform, seasonal) - **Lead Time Variability**: Random delays in shipments (min/max/mean) - **Supplier Reliability**: Partial fulfillment, quality issues,

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

disruptions - **Capacity Constraints**: Production limits, storage limits, transportation capacity - **Market Dynamics**: Price elasticity, competitive pressures, seasonal trends

Example Configurations: - **Default TBG**: Retailer → Wholesaler → Distributor → Factory (4 sites, 3 transportation lanes) - **Three FG TBG**: 3 finished goods, 3 components, 2 packaging, 1 raw material (9 sites, 12 transportation lanes) - **Complex SC**: 20+ sites with convergent manufacturing and multi-channel distribution

Planning Hierarchies (AWS Supply Chain Aligned)

Three-Dimensional Hierarchy System:

The platform supports AWS Supply Chain-aligned planning hierarchies across three dimensions, enabling planning at different levels of aggregation:

1. Site/Geographic Hierarchy:

```
Company (ACME Corp)
├── Region (Americas, EMEA, APAC)
│   ├── Country (USA, Canada, Mexico)
│   │   ├── State (California, Texas, New York)
│   │   │   └── Site (Los Angeles DC, Dallas Factory)
```

2. Product Hierarchy:

```
Category (Beverages)
├── Family (Beer)
│   ├── Group (Craft Beer)
│   │   └── Product/SKU (IPA 6-pack 12oz)
```

3. Time Bucket Hierarchy:

```
Year (Strategic)
├── Quarter (Network Design)
│   ├── Month (S&OP)
│   │   ├── Week (MPS)
│   │   │   ├── Day (MRP)
│   │   │   │   └── Hour (ATP/Execution)
```

Planning Type Configuration by Hierarchy Level:

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Planning Type	Site Level	Product Level	Time Bucket	Horizon	Powell Class	GNN Model
Execution	Site	SKU	Hour	1 week	VFA	Execution tGNN
MRP	Site	SKU	Day	13 weeks	VFA	Execution tGNN
MPS	Site	Group	Week	6 months	CFA	Hybrid
S&OP	Country	Family	Month	24 months	CFA	S&OP GraphSAGE
Capacity	Site	Group	Month	18 months	CFA	S&OP GraphSAGE
Strategic	Region	Category	Quarter	5 years	DLA	S&OP GraphSAGE

Hierarchical Consistency (Powell Framework): - Higher levels (S&OP) compute policy parameters θ via CFA - Lower levels (Execution) make decisions Q(s,a) via VFA respecting θ - Consistency constraint: V_execution ≈ E[V_tactical | policy(θ)] - Tolerance: Typically <10% deviation from parent plan

Group Administrator Configuration: - Each hierarchy configuration is customizable per group (organization) - Administrators select hierarchy levels for each planning type - Configure planning horizons, frozen periods, and update frequencies - Set Powell policy class and GNN model type per planning activity

Synthetic Data Generation (AI-Guided Setup Wizard)

Purpose: Enable rapid deployment and testing by generating realistic synthetic supply chain data for new organizations. A Claude-powered wizard guides system administrators through creating complete, archetype-based configurations.

Three Company Archetypes:

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Archetype	Description	Network Structure	Primary KPIs	Agent Mode
Retailer	Multi-channel retail operations	CDCs → RDCs → Stores + Online	Fill Rate, Inventory Turns, DOS	Copilot
Distributor	Wholesale distribution	NDCs → RDCs → LDCs + Kitting	OTIF, Order Fill Rate, Cycle Time	Copilot
Manufacturer	Production-focused operations	Plants → Sub-Assembly → Component → FG DCs	Gross Margin, Production Efficiency	Autonomous

AI-Guided Wizard Flow:

```
1. Welcome & Archetype Selection → Choose company type (Retailer/Distributor/Manufacturer)
2. Company Details → Name, group, admin credentials
3. Network Configuration → Sites, suppliers, customers (archetype defaults provided)
4. Product Configuration → SKUs, categories, families
5. Demand Configuration → Pattern (seasonal/trending/promotional), seasonality amplitude
6. Agent Configuration → Mode (none/copilot/autonomous), enable GNN/LLM/TRM
7. Review & Generate → Final confirmation, create all entities
```

What Gets Generated: - **Organization**: Group (company) and administrator user with default credentials - **Supply Chain Network**: Sites (DCs, plants, stores), transportation lanes (transportation links), products (SKUs) - **Hierarchies**: Site hierarchy (Company→Region→Country→Site), Product hierarchy (Category→Family→Group→Product) - **Planning Data**: Forecasts with P10/P50/P90 percentiles, inventory policies (DOC-based safety stock) - **Planning Configs**: MPS, MRP, S&OP configurations with Powell policy class assignments - **AI Agents**: Agent configurations per archetype with recommended strategies

Archetype-Specific Defaults:

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Parameter	Retailer	Distributor	Manufacturer
Sites	61 (2 CDC, 6 RDC, 50 stores, 3 online)	34 (2 NDC, 8 RDC, 20 LDC, 4 kitting)	31 (3 plants, 6 sub-assy, 8 comp, 14 DCs)
Suppliers	10	15	40 (25 raw + 15 tier-1)
Products	200 (5 cat × 4 fam × 10 SKU)	720 (8 cat × 6 fam × 15 SKU)	160 (4 cat × 5 fam × 8 SKU)
Safety Stock	14 days	10 days	7 days
Service Level	95%	97%	93%
Demand Pattern	Seasonal (30% amplitude)	Trending (2%/month)	Promotional (spikes)

Aggregation/Disaggregation Services (Powell Framework):

The platform implements Powell's framework insight that hierarchical planning requires both **state abstraction** (aggregation) and **policy-based allocation** (disaggregation):

Aggregation Service - State Abstraction: - **Powell Principle**: Reduce dimensionality for tractable optimization at higher levels - **Methods**: SUM, AVERAGE, WEIGHTED_AVERAGE, MIN, MAX, COUNT, VARIANCE, PERCENTILE - **Use Case**: Roll up site-level inventory to country level for S&OP planning - **Implementation**: `backend/app/services/aggregation_service.py`

Disaggregation Service - Policy-Based Allocation: - **Powell Principle**: Disaggregation proportions are a **policy decision**, not a fixed transformation - **Three Powell Policy Classes Supported**:

Method	Powell Class	Description
PROPORTIONAL	PFA	Use historical proportions (simple, stable)
LEARNED	CFA	Optimize proportions from Plan vs. Actual data
VALUE_BASED	VFA	Allocate to maximize downstream value function
CAPACITY_WEIGHTED	CFA	Weight by available capacity at each site
FORECAST_DRIVEN	CFA	Weight by forecasted demand at each site

- **Key Insight**: The `LEARNED` method trains on historical data to find optimal splits: `θ_split = argmin Σ |actual_proportion - predicted_proportion|²`
- **Use Case**: Distribute monthly S&OP family plan to weekly site-SKU MPS
- **Implementation**: `backend/app/services/disaggregation_service.py`

Hierarchical Consistency Enforcement:

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

```
V_aggregated(S_country,family,month) ≈ E[V_detailed(S_site,sku,week) | disaggregate]
```

- Default tolerance: 10% deviation between parent and child plans
- Configurable per planning type via `consistency_tolerance` parameter

API Endpoints: - `POST /api/v1/synthetic-data/wizard/sessions` - Start wizard session - `POST /api/v1/synthetic-data/wizard/sessions/{id}/messages` - Send message to wizard - `POST /api/v1/synthetic-data/wizard/sessions/{id}/generate` - Generate data - `POST /api/v1/synthetic-data/generate` - Direct generation (no wizard) - `GET /api/v1/synthetic-data/archetypes` - List archetype information

Access: System Administrator only (`/admin/synthetic-data`)

AI/ML Engine Stack (Suggested Actions Generation)

The platform's three primary AI agents continuously analyze supply chain state and generate **suggested actions** for both scenario and production environments. Each agent brings unique strengths:

1. Two-Tier Graph Neural Network (GNN) - Pattern Recognition Expert: - **Strength**: Captures complex dependencies and information flow across supply chain network over time - **Powell Framework Alignment**: S&OP=CFA (policy parameters θ), Execution=VFA (decisions Q(s,a)) - **Two-Tier Architecture** (S&OP + Execution):

S&OP GraphSAGE (Medium-Term / Strategic Planning): - **Purpose**: Network structure analysis, risk scoring, bottleneck detection - **Update Frequency**: Weekly/Monthly or on topology changes - **Architecture**: GraphSAGE with neighbor sampling, optimized for 50+ node networks - **Outputs**: Criticality scores, concentration risk, resilience scores, safety stock positioning multipliers - **Scalability**: O(edges) complexity vs $O(n^2)$ for attention, handles large supply chains efficiently

Execution tGNN (Short-Term / Operational): - **Purpose**: Real-time order decisions, demand sensing, exception detection - **Update Frequency**: Daily/Real-time - **Architecture**: Temporal GNN (GAT + GRU) consuming S&OP structural embeddings - **Inputs**: S&OP embeddings + transactional data (orders, shipments, inventory) - **Outputs**: Order recommendations, demand forecasts, exception probability, propagation impact

Shared Foundation: - S&OP structural embeddings are cached and fed to Execution model - Structural context (slow-changing) + temporal dynamics (fast-changing) - `HybridPlanningModel` provides unified interface for both tiers

- **Performance**: 85-92% accuracy on demand prediction, 15-30% cost reduction vs. naive policies

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

- **Best For**: Complex multi-echelon networks with long-term temporal patterns

2. LLM Multi-Agent System (GPT-4) - Strategic Reasoning Expert: - **Strength**: Natural language reasoning, strategic planning, contextual decision-making - **Site Agents**: GPT-4-based agents for each supply chain role (retailer, wholesaler, distributor, factory) - **Supervisor Agent**: Reviews and improves site agent decisions with global context - **Global Planner**: Coordinates network-wide strategy and optimization - **Tool Registry**: JSON schemas for structured decision-making - **Explainability**: Natural language explanations for every recommendation - **Fallback**: Heuristic policies when LLM unavailable - **Best For**: Complex scenarios requiring strategic reasoning, trade-off analysis

3. Tiny Recursive Model (TRM) - Speed and Efficiency Expert: - **Strength**: Ultra-fast inference (<10ms per decision), low computational cost - **Powell Classification**: VFA (Value Function Approximation) - fast policy execution - **Architecture**: 7M parameter transformer with recursive refinement (3 iterations) - **Performance**: 90-95% accuracy vs. optimal policies, 100+ decisions/second - **Best For**: Real-time decision-making, high-volume planning, edge deployment

Multi-Agent Consensus for Suggested Actions: - **Weighted Ensemble**: Combine agent recommendations using learned weights (e.g., LLM: 45%, GNN: 38%, TRM: 17%) - **Confidence Scoring**: Agreement between agents indicates recommendation reliability - **Adaptive Learning**: Weights automatically adjust based on observed performance (5 learning algorithms: EMA, UCB, Thompson Sampling, Performance-based, Gradient Descent) - **Context-Agnostic**: Same agents and consensus logic work for scenarios and production - **Transfer Learning**: Weights learned in scenarios transfer to production deployment

3. Reinforcement Learning (RL): - Algorithms: Proximal Policy Optimization (PPO), Soft Actor-Critic (SAC) - Reward Function: Minimize total supply chain cost + service level penalties - Training Environment: SimPy-based discrete event simulation - Exploration: Curriculum learning from simple to complex supply chains

4. Classical Optimization: - Mixed-Integer Linear Programming (MILP) for base-stock policies - Dynamic programming for multi-stage inventory optimization - Simulation-optimization for complex scenarios

5. Tiny Recursive Model (TRM): - Architecture: 7M parameter transformer with recursive refinement (3 iterations) - Input: Per-node state (inventory, backlog, pipeline, demand history, role, position) - Output: Optimal order quantity with <10ms inference time - Training: 5-phase curriculum learning from simple to complex supply chains - Performance: 90-95% accuracy vs optimal policies, 20-35% cost reduction vs naive agents

---

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

AI Model Training Workflows

TRM (Tiny Recursive Model) Training Pipeline

Philosophy: Progressive curriculum learning from simple to complex scenarios with optimal policy imitation.

Training Architecture

Model Specifications: - **Parameters**: 7 million (compact, fast) - **Architecture**: 2-layer transformer + recursive refinement - **Embedding Dimension**: 512 (d_model) - **Attention Heads**: 8 (multi-head attention) - **Refinement Steps**: 3 (iterative reasoning) - **Decision Head**: Order quantity prediction - **Value Head**: State value estimation (for RL-style training)

5-Phase Curriculum Learning

Phase 1: Single-Site Base Stock (Simplest) - **Scenario**: Solo inventory management - **Topology**: 1 site (no upstream/downstream) - **Policy**: Optimal base stock (provably optimal) - **Dataset**: 10,000 samples - **Training Time**: ~30 minutes (GPU) - **Learning**: Basic inventory control

Phase 2: 2-Site Supply Chain - **Scenario**: Simple retailer-wholesaler chain - **Topology**: 2 sites (linear chain) - **Policy**: Coordinated base stock - **Dataset**: 10,000 samples - **Training Time**: ~30 minutes (GPU) - **Learning**: Order propagation, basic bullwhip

Phase 3: 4-Site Beer Game - **Scenario**: Classic Beer Game - **Topology**: 4 sites (Retailer → Wholesaler → Distributor → Factory) - **Policy**: Tuned PID controller - **Dataset**: 10,000 samples - **Training Time**: ~30 minutes (GPU) - **Learning**: Full supply chain dynamics, amplification effects

Phase 4: Multi-Echelon Variations - **Scenario**: Complex topologies - **Topology**: 3-6 sites (varied structures) - **Policy**: Adaptive PID with forecast - **Dataset**: 10,000 samples - **Training Time**: ~45 minutes (GPU) - **Learning**: Generalization to different structures

Phase 5: Production Scenarios - **Scenario**: Real-world constraints - **Topology**: Manufacturing with BOMs, capacity limits - **Policy**: Advanced optimization - **Dataset**: 10,000 samples - **Training Time**: ~45 minutes (GPU) - **Learning**: Manufacturing constraints, multi-product coordination

Total Training Time: ~2.5 hours on NVIDIA GPU, ~8-12 hours on CPU

Data Generation Process

Generator: trm_curriculum_generator.py

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Per-Sample Structure:

```
{
    'inventory': float,              # Current inventory level
    'backlog': float,                # Current backlog
    'pipeline': List[float],         # Incoming shipments (lead time window)
    'demand_history': List[float],   # Recent demand (7-14 periods)
    'node_type': str,                # Role (retailer, wholesaler, etc.)
    'node_position': int,            # Position in chain (0=downstream)
    'target_order': float,           # LABEL: Optimal order quantity
    'target_value': float            # LABEL: State value (expected cost)
}
```

Demand Patterns: - Random: Normal distribution with volatility - Seasonal: Cyclic patterns (weekly/monthly) - Step: Sudden demand shifts - Trend: Linear growth/decline

Label Generation: - Phases 1-2: Optimal base stock policy (provably optimal) - Phases 3-5: Tuned PID controller (near-optimal heuristic)

Training Workflow

Step 1: Data Generation

```
# Generate Phase 1 dataset
python -m app.simulation.trm_curriculum_generator \
    --phase 1 \
    --num-samples 10000 \
    --output data/trm/phase1_dataset.pt
```

Step 2: Model Training

```
# Train Phase 1
python scripts/training/train_trm.py \
    --phase 1 \
    --epochs 10 \
    --device cuda \
    --batch-size 32 \
    --learning-rate 1e-4
```

Step 3: Validation & Checkpointing - Validates on 20% held-out test set - Saves checkpoint: `checkpoints/trm/trm_phase1_epoch10.pt` - Logs metrics: train_loss, val_loss, learning_rate

Step 4: Progressive Loading - Load Phase N-1 checkpoint - Continue training on Phase N data - Fine-tune existing knowledge

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Step 5: Inference Deployment - Load final checkpoint - Model automatically used by TRM agents - Fallback to base stock heuristic if model unavailable

Training via UI

TRM Dashboard (http://localhost:8088/admin/trm):

1. **Training Tab**:
2. Select curriculum phase (1-5)
3. Configure epochs, batch size, learning rate
4. Advanced settings: d_model, attention heads, layers
5. Click "Start Training"
6. **Real-Time Monitoring**:
7. Live loss charts (train/validation)
8. Epoch progress bar
9. Training status updates (every 2 seconds)
10. Automatic checkpoint saving
11. **Model Manager Tab**:
12. View available checkpoints
13. Load trained model
14. Select device (CPU/CUDA)
15. Unload model
16. **Testing Tab**:
17. Test model with custom inputs
18. Predefined scenarios (stable, spike, drop, high backlog)
19. View predicted order quantities
20. Validate model performance

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Training via CLI

```
 # Quick Phase 1 training (30 min)
make train-trm TRM_PHASE=1 TRAIN_EPOCHS=10

# Full curriculum training (2.5 hours)
cd backend
python scripts/training/train_trm.py \
    --phase 5 \
    --epochs 10 \
    --device cuda \
    --batch-size 32 \
    --learning-rate 1e-4 \
    --num-samples 10000

# Custom training with resume
python scripts/training/train_trm.py \
    --phase 3 \
    --epochs 20 \
    --device cuda \
    --resume checkpoints/trm/trm_phase2_epoch10.pt \
    --batch-size 64 \
    --learning-rate 5e-5
```

Loss Function

Multi-Objective Loss:

```
 action_loss = MSE(predicted_orders, target_orders)
value_loss = MSE(predicted_values, target_values)
total_loss = action_loss + 0.5 * value_loss
```

Optimization: - Optimizer: AdamW (weight decay = 1e-5) - Gradient clipping: max_norm = 1.0 - LR scheduler: ReduceLROnPlateau (patience=5) - Validation split: 80/20 train/test

Output & Checkpoints

Checkpoint Directory: `backend/checkpoints/trm/`

Files:

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

```
 trm_phase1_epoch10_20260117.pt  (~28MB)
trm_phase2_epoch10_20260117.pt  (~28MB)
trm_phase3_epoch10_20260117.pt  (~28MB)
trm_phase4_epoch10_20260117.pt  (~28MB)
trm_phase5_epoch10_20260117.pt  (~28MB)
```

Each checkpoint contains: - Model state dict - Optimizer state - Training configuration (hyperparameters) - Training history (loss curves) - Metadata (timestamp, phase, epochs)

Performance Metrics

Inference Speed: <10ms per decision (100+ decisions/second) **Model Size**: ~28MB on disk, ~100MB in RAM **Accuracy**: 90-95% vs optimal policies **Cost Reduction**: 20-35% vs naive agents **Generalization**: Works on unseen topologies (within training distribution)

GNN (Graph Neural Network) Training Pipeline

Philosophy: Learn supply chain dynamics from simulated scenario trajectories using graph message passing.

Training Architecture

Model Specifications: - **Parameters**: 128 million+ (heavy, expressive) - **Architecture**: Graph Attention Network (GAT) + Temporal Convolutional Network (TCN) - **Node Embedding**: 256 dimensions - **Attention Heads**: 8 (multi-head GAT) - **Temporal Layers**: 4 (TCN for time series) - **Message Passing**: 3 rounds (neighborhood aggregation) - **Output**: Per-node demand predictions + order recommendations

Data Generation Process

Generator: generate_simpy_dataset.py using data_generator.py

SimPy-Based Beer Game Simulation:

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

```
def simulate_beer_game(T=64, agent_strategy='naive'):
    """
    Run full Beer Game simulation with 4 nodes.

    Returns per-role time series:
        - inventory: [50, 48, 52, ...]
        - backlog: [0, 2, 0, ...]
        - placed_order: [52, 50, 55, ...]
    """
```

Graph-Structured Data:

```
{
    'X': node_features,        # (num_samples, 4, 12)
    'Y': action_labels,        # (num_samples, 4, 1)
    'A_ship': adjacency_ship,  # (4, 4) - shipment edges
    'A_order': adjacency_order # (4, 4) - order edges
}
```

Node Features (12 dimensions per node): 1. Inventory level 2. Backlog level 3. Incoming orders (from downstream) 4. Incoming shipments (from upstream) 5. On-order quantity (pipeline) 6-9. One-hot role encoding (retailer, wholesaler, distributor, factory) 10. Order lead time 11. Supply lead time 12. (Reserved)

Edge Features (Adjacency Matrices): - **Shipment Adjacency**: Who ships to whom (material flow) - **Order Adjacency**: Who orders from whom (information flow)

Training Workflow

Step 1: Dataset Generation

```
# Generate SimPy training dataset
cd backend
python scripts/training/generate_simpy_dataset.py \
    --config-name "Default TBG" \
    --num-runs 128 \
    --timesteps 64 \
    --agent-strategy naive \
    --output data/gnn/default_tbg_dataset.pt
```

Parameters: - `num_runs`: Number of scenario simulations (default: 128) - `timesteps`: Steps per simulation (default: 64) - `window`: History window for sequences (default: 52) - `horizon`: Forecast horizon (default: 1) - `agent_strategy`: naive, pid_heuristic, or llm

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Step 2: Sliding Window Extraction - Extract temporal windows from full scenario trajectories - Each window: past 52 timesteps → predict next 1 timestep - Creates ~8,000-16,000 training samples from 128 runs

Step 3: Model Training

```
# Train GNN on GPU
python scripts/training/train_gnn.py \
    --dataset data/gnn/default_tbg_dataset.pt \
    --epochs 50 \
    --device cuda \
    --batch-size 16 \
    --learning-rate 1e-4
```

Step 4: Validation & Checkpointing - Validates on 20% held-out test set - Saves checkpoint: `checkpoints/gnn/gnn_epoch50_20260117.pt` - Logs metrics: train_loss, val_loss, MAE, RMSE

Step 5: Deployment - Load checkpoint into GNN agent - Use for inference in ml_forecast agent strategy - Real-time demand prediction during scenarios

Training via CLI

```
# Quick training (generate data + train)
make train-gnn

# Train on GPU with custom parameters
make train-default-gpu \
    TRAIN_EPOCHS=50 \
    TRAIN_DEVICE=cuda \
    SIMPY_NUM_RUNS=256 \
    SIMPY_TIMESTEPS=64

# Generate data only (no training)
make generate-simpy-data \
    CONFIG_NAME="Default TBG" \
    SIMPY_NUM_RUNS=128

# Remote training (on remote GPU server)
make remote-train \
    REMOTE=user@gpu-server \
    EPOCHS=50 \
    DEVICE=cuda
```

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Training Parameters

Exposed in Makefile: - `CONFIG_NAME`: Supply chain config (default: "Default TBG") - `SIMPY_NUM_RUNS`: Simulation runs (default: 128) - `SIMPY_TIMESTEPS`: Steps per run (default: 64) - `SIMPY_WINDOW`: History window (default: 52) - `SIMPY_HORIZON`: Forecast horizon (default: 1) - `TRAIN_EPOCHS`: Training epochs (default: 10) - `TRAIN_DEVICE`: cuda or cpu (default: cuda)

Code-Only (in train_gnn.py): - Hidden dimensions: [256, 128, 64] - GAT attention heads: 8 - TCN kernel size: 3 - Dropout: 0.3 - Weight decay: 1e-5 - Learning rate decay: 0.95 per epoch

Graph Message Passing

Forward Pass: 1. **Node Embedding**: Encode 12 features → 256-dim embedding 2. **GAT Layer 1**: Aggregate neighbor features with attention 3. **GAT Layer 2**: Refine embeddings with multi-head attention 4. **GAT Layer 3**: Final neighborhood aggregation 5. **Temporal Processing**: TCN across time dimension 6. **Output Layer**: Predict demand + optimal orders per node

Attention Mechanism:

```
 # Compute attention weights between nodes
attention_weights = softmax(
    LeakyReLU(
        W * [node_i_features || node_j_features]
    )
)

# Aggregate neighbor features
aggregated = sum(attention_weights * neighbor_features)
```

Loss Function

Multi-Task Loss:

```
 demand_loss = MSE(predicted_demand, actual_demand)
action_loss = MSE(predicted_orders, agent_orders)
total_loss = 0.7 * demand_loss + 0.3 * action_loss
```

Optimization: - Optimizer: Adam (no weight decay initially) - Learning rate: 1e-4 → 1e-6 (exponential decay) - Batch size: 16 (graph batching) - Validation frequency: Every epoch

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Output & Checkpoints

Checkpoint Directory: `backend/checkpoints/gnn/`

Files:

```
 gnn_default_tbg_epoch50_20260117.pt     (~500MB)
gnn_three_fg_epoch50_20260117.pt        (~500MB)
gnn_complex_sc_epoch100_20260117.pt     (~500MB)
```

Each checkpoint contains: - Model state dict - Optimizer state - Graph structure (adjacency matrices) - Training configuration - Performance metrics (MAE, RMSE, R^2)

Performance Metrics

Inference Speed: ~50-100ms per graph forward pass **Model Size**: ~500MB on disk, ~2GB in RAM (GPU) **Demand Prediction Accuracy**: 85-92% (within 15% of actual) **Cost Reduction**: 15-30% vs naive agents **Bullwhip Reduction**: 20-40% variance reduction upstream

Training Time Comparison

Dataset Size	CPU Time	GPU Time
128 runs	~6 hours	~1.5 hours
256 runs	~12 hours	~3 hours
512 runs	~24 hours	~6 hours

Recommendation: Use GPU for training (5-8x speedup)

Training Data Comparison: TRM vs GNN

See detailed comparison: TRM_VS_GNN_TRAINING_DATA.md

Quick Summary:

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Aspect	TRM	GNN
Data Format	Per-node flat tensors	Graph-structured tensors
Topology	Variable (1-N nodes)	Fixed (4 nodes)
Labels	Optimal policies	Agent trajectories
Curriculum	5 progressive phases	Single-pass
Dataset Size	50K samples	8K-16K windows
Generation Time	~30 min	~1-2 hours
Training Time	~2.5 hours (GPU)	~4-6 hours (GPU)
Inference	<10ms	~50-100ms
Flexibility	Generalizes to varied topologies	Beer Game specific

Key Insight: TRM and GNN use **completely different data** and cannot share datasets without conversion.

Technology Stack

Backend: - Python 3.10+, FastAPI (async REST API) - SQLAlchemy 2.0 (async ORM) - PyTorch 2.2 + PyTorch Geometric (GNN) - Stable-Baselines3 (RL) - OpenAI API (LLM agents) - SimPy (discrete event simulation)

Frontend: - React 18 (functional components, hooks) - Material-UI 5 (enterprise UI) - Recharts + D3-Sankey (visualizations) - WebSocket (real-time updates)

Infrastructure: - Docker + Docker Compose (containerization) - PostgreSQL 16 (relational database) - Redis (caching, session management) - Nginx (reverse proxy, load balancing) - GPU support (NVIDIA CUDA for ML training)

SAP S/4HANA & IBP Integration

Overview

The Continuous Autonomous Planning Platform provides enterprise-grade integration with SAP S/4HANA and SAP Integrated Business Planning (IBP) for seamless data exchange, AI-powered planning enhancement, and real-world validation of supply chain strategies.

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Integration Architecture

External System Integration Strategy:

```
┌──────────────────┐      ┌──────────────────────────┐      ┌──────────────────┐
│   Databricks     │──────│  Autonomy Platform       │──────│  External        │
│   (Data Lake)    │      │  (Supply Chain Exec)     │      │  Demand Planning │
└──────────────────┘      └──────────────────────────┘      └──────────────────┘

        ↓                           ↓                              ↓
   Master Data              Execution & Visibility         Forecast Generation
   ERP Integration          Order Management                Consensus Planning
   Data Transformation      ATP/CTP, Inventory              ML Forecasting
   Analytics                Agent-Based Planning            Promotional Planning
```

Architectural Decisions: 1. **Data Lake**: Databricks (external platform) - ERP connector setup via Databricks - Data mapping and transformation in Databricks - Data quality validation pipeline - Future: Databricks connector for Autonomy analytics

1. **Demand Planning**: External system (API integration)
2. Forecast generation in external system
3. Autonomy consumes via REST API (read-only)
4. Display current demand plan + delta analysis
5. No forecast modification in Autonomy
6. **Supply Planning & Execution**: Autonomy (core platform)
7. AWS SC-compliant data model
8. MPS/MRP planning with AI agents
9. Order management (PO/TO/MO)
10. ATP/CTP and inventory optimization
11. Real-time execution and collaboration

Bidirectional Data Flow:

```
SAP S/4HANA ↔ Databricks ↔ Autonomy Platform ↔ External Demand Planning
     ↓            ↓              ↓                        ↓
   Master Data  Data Lake   Simulations/Execution    Forecasts
   Actuals      Analytics   Simulation/Training      Consensus
   Orders       Transform   AI Agents                Planning
```

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Data Ingestion Methods

1. Direct SAP Integration (Real-Time)

SAP OData API Integration: - RESTful API connectivity to SAP S/4HANA OData services - Authentication via OAuth 2.0 or SAP Cloud Platform - Real-time data pull from SAP tables (MARA, MARC, EBAN, EKKO, LIKP, etc.) - Support for SAP BAPIs for transactional operations

Supported Data Entities: - **Master Data**: Materials (MARA), Plants (T001W), Vendors (LFA1), Customers (KNA1) - **Inventory**: Stock levels (MARD), reservations (RESB), goods movements (MKPF) - **Purchasing**: Purchase orders (EKKO/EKPO), requisitions (EBAN) - **Sales**: Sales orders (VBAK/VBAP), deliveries (LIKP/LIPS) - **Production**: Planned orders (PLAF), production orders (AFKO) - **Demand**: Consumption data, forecast (FCST), demand planning (IBP)

SAP HANA Direct Query (for S/4HANA on HANA): - Direct SQL queries to SAP HANA database views - CDS (Core Data Services) view consumption - Real-time analytics on SAP data without extraction

SAP IBP API Integration: - Connect to IBP Planning Area APIs - Import demand forecasts, supply plans, inventory targets - Export AI-generated recommendations back to IBP - Support for IBP Key Figures and Time Series data

2. Flat File Integration (Batch)

Supported Formats: - **CSV**: Comma-separated values with configurable delimiters - **Excel**: .xlsx files with multi-sheet support - **JSON**: Structured hierarchical data - **XML**: SAP IDoc format, custom schemas - **Parquet**: Columnar format for large datasets

File Transfer Methods: - SFTP/FTP upload to designated folders - AWS S3 / Azure Blob Storage integration - Direct file upload via UI - Scheduled batch imports (cron jobs)

Standard File Templates:

```
 # inventory.csv
material_id,plant,storage_location,quantity,unit,timestamp
MAT001,1000,0001,1500,EA,2026-01-17T08:00:00Z

# demand.csv
material_id,location,date,quantity,source
MAT001,DC-EAST,2026-01-17,250,actual

# supply_plan.csv
material_id,vendor,plant,planned_date,quantity,lead_time_days
MAT002,V001,1000,2026-01-24,1000,7
```

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Data Mapping Configuration: - Field mapping UI to align SAP field names to Autonomy schema - Transformation rules (unit conversions, date formats, aggregations) - Data validation and cleansing rules - Error handling and logging

Claude AI Data Curation

Automated Data Cleansing: - **Missing Data Imputation**: Claude analyzes patterns and fills gaps intelligently - **Anomaly Detection**: Identifies outliers, data quality issues, inconsistencies - **Data Standardization**: Harmonizes units, date formats, naming conventions - **Duplicate Detection**: Finds and merges duplicate records across systems

Semantic Data Enrichment: - **Natural Language Processing**: Extract insights from free-text fields (e.g., vendor notes, order comments) - **Contextual Understanding**: Claude infers relationships between entities (e.g., which materials are substitutes) - **Classification**: Auto-categorize materials by ABC class, demand patterns, criticality - **Entity Resolution**: Match SAP material codes to Autonomy products with fuzzy matching

Data Transformation Workflows:

```
SAP Raw Data → Claude Curation → Validated Data → Autonomy Schema
     ↓               ↓                  ↓                 ↓
   Tables      Cleansing          Mapping         Supply Chain
   Views       Enrichment         Validation      Configuration
   Files       Inference          Approval        Scenarios
```

Claude Curation Features: 1. **Intelligent Field Mapping**: - Claude suggests optimal field mappings based on field names and data patterns - Learns from user corrections to improve future mappings - Handles SAP custom fields (ZFIELD*) and custom tables

1. **Data Quality Scoring**:
2. Assigns quality scores to incoming data (0-100%)
3. Flags low-quality records for review
4. Provides explanations for quality issues
5. **Automated Documentation**:
6. Claude generates data lineage documentation
7. Creates data dictionary entries automatically
8. Explains transformation logic in natural language
9. **Conversational Curation**:
10. Users can ask Claude: "Why is inventory for MAT001 flagged?"
11. Claude explains data issues and suggests fixes

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

12. Interactive approval/rejection of curation suggestions

SAP Data Management UI

The platform includes a comprehensive SAP Data Management interface accessible from Administration > SAP Data Management:

Connection Management: - Configure connections to SAP S/4HANA, APO, ECC, or BW systems - Support for multiple connection methods: RFC, CSV file import, OData API, IDoc - Connection testing and validation with status tracking - Manage multiple connections per group (dev/test/prod)

Z-Table/Z-Field Handling: - Automatic discovery of Z-tables (custom SAP tables) - AI-powered fuzzy matching for Z-field to AWS SC entity mapping - Pattern-based recognition of common SAP naming conventions - Confidence scoring (High/Medium/Low) for each suggested mapping - User confirmation workflow to learn from corrections

Field Mapping Features: | Mapping Method | Description | Confidence | |---|---|---|---| | Exact Match | Direct field name match | 95%+ | | Pattern Match | SAP naming convention patterns (MATNR→product_id) | 90%+ | | Fuzzy Match | Levenshtein + token similarity | 70-90% | | AI Suggested | Claude analysis for ambiguous fields | 50-75% | | Learned | From previous user confirmations | 98% |

Data Ingestion Monitoring: - Real-time job tracking with progress indicators - Data quality scoring across 5 dimensions (Completeness, Accuracy, Consistency, Uniqueness, Validity) - Anomaly detection for unexpected data patterns - Trend analysis showing quality changes over time

Insights & Actions Dashboard: - AI-generated insights categorized by severity (Critical, Error, Warning, Info) - Suggested remediation actions with one-click execution - Action workflow (Suggested → In Progress → Completed) - Integration with existing exception workflows

Write-Back to SAP

1. Recommendations Export

AI-Generated Plans to SAP IBP: - Export TRM/GNN-optimized order quantities to IBP Planning Areas - Push demand forecasts back to IBP Demand module - Update safety stock recommendations in IBP Supply Planning - Write inventory targets to IBP Key Figures

Export Format:

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

```
{
  "planning_area": "DEMAND_FORECAST",
  "key_figures": [
    {
      "material": "MAT001",
      "location": "DC-EAST",
      "date": "2026-01-24",
      "forecast_quantity": 275.5,
      "confidence": 0.92,
      "model": "TRM",
      "rationale": "15% seasonal uplift, 5% trend"
    }
  ]
}
```

2. Transactional Write-Back

Create Purchase Requisitions in S/4HANA: - AI agents generate recommended orders - Human approval workflow - Automatic PR creation via SAP BAPI (BAPI_PR_CREATE) - Track PR status in Autonomy platform

Update Inventory Targets: - Write optimized reorder points to material master (MARC-MINBE) - Update safety stock levels (MARC-EISBE) - Modify service levels by material/plant

Scenario Planning Export: - Export multiple "what-if" scenarios from Autonomy simulations - Load scenarios into SAP IBP Scenario Comparison - Enable side-by-side analysis of AI vs. manual plans

3. Audit Trail & Governance

Complete Traceability: - Every write-back operation logged with timestamp, user, rationale - Claude generates audit summaries in natural language - Compliance with SOX, GDPR, data residency requirements - Rollback capability for all SAP updates

Approval Workflows: - Multi-level approvals before SAP write-back - Role-based access control (planner, manager, director) - Email/Slack notifications for pending approvals - Automated approval for low-risk changes (configurable thresholds)

Integration Capabilities

Real-Time Synchronization

Incremental Data Sync: - CDC (Change Data Capture) from SAP tables - WebSocket push notifications on SAP data changes - Beer Game platform reacts to real-world supply chain events - Sub-second latency for critical updates

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Event-Driven Architecture: - SAP Business Events trigger supply chain simulations - Examples: PO delayed → run risk simulation, Demand spike → AI re-plan - Event bus integration (Kafka, RabbitMQ, AWS EventBridge)

Hybrid Planning Workflows

SAP IBP + AI Hybrid: 1. SAP IBP generates statistical forecast (baseline) 2. Autonomy TRM/GNN models generate AI forecast (enhanced) 3. Claude compares forecasts, explains differences 4. Planner chooses: SAP forecast, AI forecast, or weighted blend 5. Final forecast written back to IBP

Consensus Planning: - Autonomy simulations validate SAP plans under stress scenarios - Multi-agent LLM provides "second opinion" on SAP recommendations - Identify risks in SAP plans before execution (bullwhip, shortages, excess inventory)

Data Reconciliation

SAP vs. Autonomy Comparison: - Side-by-side view of SAP actuals vs. Autonomy predictions - Highlight discrepancies and investigate root causes - Claude explains variances (e.g., "SAP shows 10% lower demand due to canceled orders")

Feedback Loop: - Platform learns from SAP actuals to improve models - Model retraining triggered by drift detection - Continuous improvement of AI accuracy

Enterprise Features

SAP Security Integration

Single Sign-On (SSO): - SAML 2.0 integration with SAP Identity Management - OAuth 2.0 with SAP Cloud Platform - Support for SAP Principal Propagation - Active Directory / LDAP integration

SAP Authorization Model: - Map SAP roles to Autonomy permissions - Respect SAP authorization objects (S_TABU_DIS, M_MATE_WRK, etc.) - Row-level security based on SAP org hierarchy

Data Residency & Compliance

On-Premise Deployment: - Deploy Autonomy Platform within SAP landscape - No data leaves corporate network - Complies with data sovereignty requirements

Cloud Deployment Options: - SAP Cloud Platform Integration (CPI) connectivity - AWS / Azure / GCP with private VPN to SAP - Data encryption in transit (TLS 1.3) and at rest (AES-256)

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Regulatory Compliance: - GDPR: Data retention policies, right to deletion - SOX: Audit trails, segregation of duties - FDA 21 CFR Part 11: Electronic signatures for pharma/medical device - GxP: Validated systems for regulated industries

Performance & Scalability

High-Volume Data Handling: - Process 1M+ SAP material master records - Ingest 10M+ daily transactions (orders, shipments, receipts) - Real-time analytics on 100M+ historical data points - Parallel processing with distributed compute

Caching & Optimization: - Redis caching for frequently accessed SAP data - Materialized views for aggregated metrics - Incremental refresh (only changed data) - Query optimization for SAP HANA

Integration Use Cases

Use Case 1: AI-Enhanced Demand Planning

Workflow: 1. **Ingest**: Pull 2 years of historical sales data from SAP (VBAK/VBAP) 2. **Curate**: Claude cleanses, identifies seasonality, fills gaps 3. **Train**: TRM model trains on SAP historical actuals 4. **Forecast**: Generate 12-month forecast with confidence intervals 5. **Compare**: Side-by-side with SAP IBP statistical forecast 6. **Validate (Digital Twin Testing)**: Run forecast through digital twin scenario simulations with stochastic demand variability—test forecast accuracy under different scenarios before production deployment 7. **Approve**: Planner reviews Claude's explanations, simulation results, and approves 8. **Write-Back**: Export validated forecast to IBP Planning Area

Result: 15-25% forecast accuracy improvement, 40% reduction in planner workload, 90%+ confidence from digital twin validation

Use Case 2: Safety Stock Optimization

Workflow: 1. **Ingest**: Pull current safety stock levels (MARC-EISBE) and service level targets 2. **Digital Twin Simulation**: Run 1,000+ Monte Carlo simulations per SKU in digital twin scenarios—test different safety stock levels under demand/lead time variability 3. **Optimize**: TRM/GNN agents analyze simulation results and find optimal safety stock for 98% service level 4. **Explain**: Claude generates rationale for each recommendation with simulation evidence 5. **Review**: Planner reviews AI suggestions with digital twin test results (service level achieved, cost impact, risk profile) 6. **Write-Back**: Update safety stock in SAP material master with validated recommendations

Result: 20-30% inventory reduction while maintaining service levels, 95%+ confidence from 1,000+ digital twin tests

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Use Case 3: Supplier Risk Assessment

Workflow: 1. **Ingest**: Pull supplier master (LFA1), PO history (EKKO), quality data 2. **Enrich**: Claude analyzes supplier performance, identifies patterns 3. **Digital Twin Testing**: Model supply disruption scenarios in digital twin scenarios—test impact of supplier failures, lead time extensions, partial fulfillment 4. **Score**: AI assigns risk scores to each supplier based on simulation results 5. **Recommend**: Suggest backup suppliers, safety stock increases, multi-sourcing strategies (all validated through digital twin testing) 6. **Alert**: Proactive warnings when high-risk suppliers have large POs, with mitigation strategies pre-tested

Result: 50% reduction in stockouts from supplier issues, validated mitigation strategies ready for immediate deployment

API Documentation

RESTful API Endpoints:

```
 POST /api/v1/sap/ingest          # Trigger SAP data import
GET  /api/v1/sap/status           # Check integration status
POST /api/v1/sap/curate           # Run Claude curation pipeline
POST /api/v1/sap/export           # Export to SAP (IBP/S4HANA)
GET  /api/v1/sap/audit-log        # Retrieve audit trail
POST /api/v1/sap/validate-mapping # Validate field mappings
GET  /api/v1/sap/data-quality     # Get data quality report
```

Webhook Support: - Register webhooks for SAP events - Receive notifications on data changes, curation completion, export success/failure

Configuration UI

SAP Connection Manager: - Configure SAP system details (host, client, credentials) - Test connectivity - Manage multiple SAP systems (dev, QA, prod)

Data Mapping Studio: - Drag-and-drop field mapping - Preview data transformations - Save mapping templates for reuse - Version control for mappings

Curation Dashboard: - View Claude curation pipeline status - Approve/reject curation suggestions - Override AI decisions with manual edits - Monitor data quality scores

Export Configuration: - Select target SAP system (IBP, S/4HANA) - Choose data entities to export - Set export frequency (real-time, hourly, daily) - Configure approval workflows

---

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Use Cases & Business Value

Use Case 1: New Hire Training

Challenge: New supply chain analysts take 6-12 months to become proficient. Traditional training is expensive and slow.

Solution: - New hires run Beer Game scenarios against AI agents - Progress through difficulty levels (Easy → Hard AI opponents) - Unlock achievements as they master concepts - Compete on leaderboards for recognition

Results: - Reduce onboarding time to 2-3 months - 85% of trainees reach proficiency benchmarks - 92% engagement scores vs. 45% for traditional training - $50K cost savings per hire (reduced training overhead)

Use Case 2: AI Agent Validation Before Deployment

Challenge: Company wants to deploy ML-based demand forecasting but stakeholders are skeptical. "How do we know it won't cause stockouts?"

Solution: - Run 100 simulated scenarios: AI agent vs. current human planner - AI wins 73% of scenarios with 22% lower total costs - Humans observe AI decisions, understand logic - Pilot deployment with human override capability

Results: - Executive buy-in achieved in 2 weeks (vs. 6-month sales cycle) - AI deployed with confidence, no override needed after 30 days - 18% reduction in inventory holding costs - 12% improvement in service levels (fewer stockouts)

Use Case 3: Supply Chain Network Redesign

Challenge: Company considering adding a regional distribution center (DC). Will it reduce costs or add complexity?

Solution: - Create two supply chain configurations (with/without DC) - Run 500 simulated scenarios for each topology - Compare total costs, service levels, bullwhip effects - Test sensitivity to demand variability, lead time changes

Results: - DC reduces costs by 14% under current demand patterns - But increases bullwhip effect by 9% - Recommendation: Add DC with visibility sharing and coordination protocols - Decision made with 95% confidence vs. "gut feeling"

Use Case 4: Performance Benchmarking

Challenge: Company suspects its ordering policies are suboptimal but lacks benchmarks.

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Solution: - Upload 12 months of historical ordering/inventory data - Platform simulates same scenarios with AI-optimal policies - Generate gap analysis report showing missed savings

Results: - Identified $2.3M annual cost reduction opportunity - 65% of excess costs from over-ordering at Distributor role - 25% from poor coordination between Wholesaler and Retailer - Implemented AI suggestions, captured 80% of savings in 6 months

Use Case 5: Collaborative Planning with Suppliers

Challenge: Retailer and supplier struggle with bullwhip effect. Lack of trust and coordination.

Solution: - Retailer and supplier run collaborative scenario together - See real-time impact of information sharing and joint planning - Negotiate visibility agreements in-scenario - Test different collaboration strategies (VMI, CPFR, etc.)

Results: - 31% reduction in demand variability at supplier - 18% reduction in inventory holding costs across both parties - Improved relationship, increased willingness to share data - Real-world collaboration agreement signed after 3 gaming sessions

---

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Competitive Advantages

vs. Enterprise Planning Systems (Kinaxis RapidResponse, SAP IBP, OMP Plus)

Feature	Kinaxis/SAP IBP/OMP	The Continuous Autonomous Planning Platform
Deployment Time	6-18 months implementation	Deploy in days with pre-built configurations
License Cost	$100K-$500K per user/year	$10K/year per seat (90% cost reduction)
Implementation Cost	$500K-$5M consulting fees	$50K-$200K one-time setup
User Training	3-6 months specialized training	2-3 weeks via gamified learning
AI/ML Capabilities	Limited, add-on modules	Native temporal GNN, LLM agents, RL optimization
Variability Modeling	Deterministic with safety stock buffers	Full stochastic simulation (demand, lead time, supplier reliability)
Human-AI Interaction	Passive acceptance of recommendations	Active gaming validates AI before production deployment
Validation	Post-implementation performance reviews	Pre-deployment testing in risk-free simulations
Learning Curve	Steep, requires consultants	Intuitive gaming interface, no coding required
Explainability	Complex optimization black boxes	Observable decisions, natural language explanations (LLM)
Adoption Risk	High—requires organizational change management	Low—simulation builds trust before deployment
Flexibility	Rigid configuration, custom dev expensive	DAG-based topology, configure any network in minutes
Integration	Requires SAP/Oracle/ERP integration	REST API, export formats, cloud or self-hosted
Vendor Lock-In	High (proprietary platforms)	Open architecture, data portability

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

vs. Traditional Supply Chain Simulation (AnyLogic, Simul8)

Feature	Traditional Simulation	The Continuous Autonomous Planning Platform
Ease of Use	Requires simulation expertise	Gaming interface, no coding required
AI Integration	Limited, custom development	7 agent strategies built-in, LLM/GNN ready
Real-Time Multiplayer	Rare, clunky	Native WebSocket support, seamless
Simulation	None	17 achievements, leaderboards, progression
Analytics	Export data, analyze elsewhere	Built-in dashboards, automated reports
Cost	$5K-$50K per license	Open source core, enterprise add-ons

vs. Supply Chain Simulation (e.g., The Fresh Connection, Beer Game Mobile Apps)

Feature	Traditional Simulations	The Continuous Autonomous Planning Platform
AI Opponents	None or scripted	7 adaptive strategies + LLM/GNN
Custom Supply Chains	Fixed scenarios	DAG editor, unlimited configurations
Real Data Integration	None	Upload historical data for analysis
Agent Training	Not possible	RL, GNN, LLM training pipelines
Enterprise Features	Consumer-focused	SSO, RBAC, audit logs, multi-tenancy
API Integration	None	REST API, WebSocket, export formats

Simulation Deep Dive: The Engine of Improvement

Achievement System Design

17 Achievements Across 5 Categories:

1. **Cost Control** (4 achievements)

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

2. "Penny Pincher": Total cost < $5,000
3. "Efficient Operator": Cost per round < $500
4. "Inventory Optimizer": Never hold > 30 units
5. "Zero Waste": Total holding cost < $1,000
6. **Service Excellence** (3 achievements)
7. "Perfect Service": 100% fill rate for 10 consecutive rounds
8. "Quick Responder": Reduce backlog by 50% in 3 rounds
9. "Zero Backlog": Complete scenario with zero stockouts
10. **Collaboration** (4 achievements)
11. "Negotiation Master": 5 successful negotiations
12. "Visibility Pioneer": Share data for 10 rounds
13. "Win-Win Participant": Both parties reduce costs after negotiation
14. "Supply Chain Diplomat": Maintain 3+ active agreements
15. **Innovation** (3 achievements)
16. "Bullwhip Tamer": Reduce demand amplification by 40%
17. "AI Adopter": Use AI suggestions for 10 rounds
18. "Steady Hand": Keep order variance < 5 units for 15 rounds
19. **Mastery** (3 achievements)
20. "Supply Chain Expert": Win 10 scenarios
21. "Leaderboard Climber": Top 10 in any leaderboard
22. "Perfect Run": Total cost < $3,000 AND 100% service level

Achievement → AI Training Mapping:

Each achievement translates to a measurable objective for AI agents:

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

```
 # Example: "Steady Hand" Achievement
# Human Goal: Keep inventory variance low
# AI Reward Function:
reward = -total_cost - 50 * variance(orders)

# Example: "Bullwhip Tamer" Achievement
# Human Goal: Reduce demand amplification
# AI Reward Function:
bullwhip_ratio = std(orders_placed) / std(orders_received)
reward = -total_cost - 100 * max(0, bullwhip_ratio - 1.2)

# Example: "Negotiation Master" Achievement
# Human Goal: Successful collaboration
# AI Reward Function (multi-agent):
reward_retailer = -cost_retailer + 20 * successful_negotiations
reward_wholesaler = -cost_wholesaler + 20 * successful_negotiations
```

Leaderboard System

6 Leaderboard Types:

1. **Lowest Total Cost**: Raw performance metric
2. **Best Service Level**: Fill rate + backlog penalties
3. **Most Improved**: Week-over-week cost reduction
4. **Achievement Hunter**: Total achievements unlocked
5. **Highest Level**: Participant progression (level = floor(sqrt(points/10)) + 1)
6. **AI Challenger**: Best performance against hardest AI

Competitive Dynamics: - Weekly leaderboard resets for fresh competition - All-time leaderboards for legacy recognition - Team leaderboards for cross-functional competitions - Anonymous peer benchmarking (see percentile rank without exposing individual data)

Social Learning: - See leaderboard leaders' replays - Learn from top performers' decision patterns - AI analyzes leader strategies, incorporates into recommendations

Progression System

Participant Levels: - Formula: `level = floor(sqrt(total_points / 10)) + 1` - Earn points for: achievements unlocked, scenarios won, cost savings, service level targets - Level 1 (Novice) → Level 10 (Expert) → Level 20 (Master) → Level 50 (Legend)

Unlockable Content: - Level 5: Access to AI suggestions - Level 10: Unlock advanced analytics dashboard - Level 15: Create custom supply chain configurations - Level 20: Train personal AI agent on own gameplay - Level 25: Access to LLM-powered agents

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Badging & Recognition: - Visual badges displayed on participant profiles - In-scenario titles (e.g., "Master of the Bullwhip", "Collaboration Champion") - Certificates for completing achievement categories - Share achievements on LinkedIn, company intranet

The Virtuous Cycle

```
┌─────────────────────────────────────────────────────────────┐
│                                                             │
│  1. Humans Run Scenarios                                    │
│     ↓                                                       │
│  2. Generate Diverse Decision Data                          │
│     ↓                                                       │
│  3. AI Learns from Human Strategies (Imitation)             │
│     ↓                                                       │
│  4. AI Competes Against Humans (Validation)                 │
│     ↓                                                       │
│  5. Humans Observe AI, Learn New Patterns                   │
│     ↓                                                       │
│  6. Humans Improve Performance (Adopt AI Insights)          │
│     ↓                                                       │
│  7. AI Learns from Improved Humans (Imitation++)            │
│     ↓                                                       │
│  8. Cycle Repeats → Continuous Improvement                  │
│     ↓                                                       │
│  [Back to Step 1]                                           │
│                                                             │
└─────────────────────────────────────────────────────────────┘
```

Measurable Outcomes: - Average participant cost reduces by 35% after 20 scenarios - AI agent performance improves by 12% per training cycle - 78% of participants report "increased confidence in AI recommendations" - 65% of participants adopt AI-suggested ordering patterns in real work

---

Implementation Roadmap

Platform Capabilities: Full AWS SC Integration

Core Supply Chain Planning (AWS SC Compliant): - Complete AWS Supply Chain feature coverage (Insights, Recommendations, Collaboration, Material Visibility, Order Planning) - ML-powered risk detection and predictive analytics - Rebalancing recommendations engine with action scoring - Real-time shipment tracking and delivery risk analytics - Team collaboration with messaging and approval workflows -

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

View-only demand planning integration with delta analysis - Comprehensive order management (PO/TO/MO/Project/Maintenance)

AI & Gaming (Production Ready): - Core Beer Game engine with multi-echelon support - 7 AI agent strategies (Naive, Conservative, Bullwhip, ML-Forecast, Optimizer, Reactive, LLM) - Temporal GNN for supply chain prediction - LLM multi-agent system (OpenAI GPT integration) - DAG-based supply chain configuration - Real-time WebSocket multiplayer gaming - Advanced analytics and bullwhip effect tracking - Simulation system (17 achievements, 6 leaderboards, progression)

Infrastructure & Integration: - Databricks data lake integration for ERP connectivity - External demand planning system integration (API-based) - Reporting and export (CSV, JSON, Excel) - Interactive tutorial and help system - Performance optimizations (database indexes, caching) - Role-based access control (RBAC) with capability-based permissions

Future Enhancements

Priority 1: Advanced AI/ML (10-15 days) - Reinforcement learning agents (PPO, SAC, A3C) - Enhanced GNN architectures (GraphSAGE, Heterogeneous GNN) - Predictive analytics dashboard for AWS SC Insights - AutoML for hyperparameter tuning - Explainable AI (SHAP values, attention visualization)

Priority 2: 3D Visualization (8-12 days) - Three.js 3D supply chain network - Geospatial mapping with real locations - Timeline animation of inventory flow - VR/AR readiness for immersive training - Performance optimization for large networks

Priority 3: Mobile Application (10-15 days) - React Native + Expo cross-platform app - Mobile-optimized scenario interface - Push notifications for scenario events - Offline mode with data sync - Mobile analytics dashboard

Priority 4: Enterprise Features (7-10 days) - SSO/LDAP integration (SAML2, OAuth2, Active Directory) - Multi-tenancy (subdomain routing, data isolation) - Advanced RBAC (granular permissions, custom roles) - Audit logging (compliance, security) - Enterprise reporting and governance

---

Business Model & Pricing

Open Core Model

Free Tier (Community Edition): - Core Beer Game engine - Default supply chain configurations - 7 AI agent strategies (excluding LLM) - Single-tenant deployment - Community support

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Professional Tier ($2,500/month): - LLM-based AI agents (GPT-4, Claude, etc.) - Custom supply chain configurations (unlimited) - Advanced analytics and reporting - Export to CSV, JSON, Excel - Email support

Enterprise Tier ($10,000/month): - SSO/LDAP integration - Multi-tenancy - Advanced RBAC and audit logging - Dedicated training and onboarding - 24/7 support with SLA - On-premise deployment option

Add-Ons: - GNN Training Service: $5,000 one-time + $500/month compute - Custom Agent Development: $15,000 per agent - Professional Services: $200/hour for consulting, custom integrations

ROI Calculator: The Continuous Autonomous Planning Platform vs. Legacy Planning Software

Example: Mid-Size Manufacturer ($500M revenue, $50M inventory)

Scenario A: Traditional Enterprise Planning (Kinaxis/SAP IBP)

Implementation Costs (Year 1): - Software licenses: 10 users × $250K = $2.5M - Implementation consulting: $2M - User training: $500K - System integration: $1M - **Total First-Year Cost: $6M**

Ongoing Costs (Annual): - Maintenance/support: $750K - Consultant retainers: $400K - User training (new hires): $200K - **Total Recurring Cost: $1.35M/year**

Time to Value: 12-18 months before operational benefits

Scenario B: The Continuous Autonomous Planning Platform

Implementation Costs (Year 1): - Platform licenses: 10 users × $10K = $100K - Setup/configuration: $75K - Training (gamified): $25K - **Total First-Year Cost: $200K**

Ongoing Costs (Annual): - Platform subscription: $100K - Support: $20K - **Total Recurring Cost: $120K/year**

Time to Value: 2-4 weeks to production deployment

Comparative Analysis

Direct Cost Savings (vs. Kinaxis/SAP IBP): - First Year: $6M - $200K = **$5.8M saved** - Year 2-5: $1.35M - $120K = **$1.23M saved per year** - 5-Year TCO: $11.4M vs. $1.08M = **$10.32M total savings (91% reduction)**

Operational Benefits (Conservative Estimates): - 10% inventory reduction: $50M × 10% = $5M freed capital - 15% inventory holding cost savings: $12.5M × 15% =

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

$1.875M/year - 20% stockout reduction: $25M × 20% = $5M revenue protection - 60% faster training: 50 analysts × $30K = $1.5M saved - **Total Annual Operational Value: $8.375M**

Total Platform ROI (First Year): - Cost Savings: $5.8M (vs. Kinaxis/SAP) - Operational Benefits: $8.375M - **Total Value: $14.175M** - **Platform Cost: $200K** - **First-Year ROI: 70.9x return**

Payback Period: 5.3 days

Key Insights: 1. **91% lower TCO** than legacy enterprise planning systems 2. **10-20x faster deployment** (weeks vs. 12-18 months) 3. **Gaming reduces training costs by 60%** vs. consultant-led workshops 4. **Risk-free AI validation** eliminates costly deployment failures 5. **No vendor lock-in**: migrate data freely, self-hosted or cloud deployment

---

Competitive Differentiation Summary

A Modern Alternative to Legacy Planning Software

Why Replace Kinaxis/SAP IBP/OMP with The Continuous Autonomous Planning Platform?

1. **90% Cost Reduction**: $100K-$500K/user/year → $10K/user/year
2. **10-20x Faster Deployment**: Weeks instead of 12-18 months
3. **AI-First Architecture**: Native temporal GNN, LLM agents, RL optimization (not bolt-on add-ons)
4. **Risk-Free Validation**: Test AI in gaming environments before production deployment
5. **Gamified Training**: 60% training cost reduction vs. consultant-led workshops
6. **Stochastic Variability**: Model real-world uncertainty (demand, lead times, supplier reliability)
7. **Transparent AI**: Observable decisions, natural language explanations, builds trust
8. **No Vendor Lock-In**: Open architecture, data portability, self-hosted or cloud
9. **Intuitive UX**: Gaming interface vs. complex ERP-style screens
10. **Continuous Learning**: AI improves from human gameplay, not just periodic retraining

The Only Platform That Combines:

1. **Gaming + Analytics + Planning** in one unified environment (Kinaxis/SAP require separate tools)

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

2. **Human + AI collaboration** with trust-building through competition (legacy systems demand blind trust)
3. **Multi-echelon variability modeling** beyond deterministic MRP/DRP calculations
4. **7+ AI strategies** including cutting-edge LLM and GNN approaches (not rigid optimization solvers)
5. **Simulation for confidence building** and continuous agent improvement (legacy training is workshop-based)
6. **Real-time multiplayer** with WebSocket and observable AI decisions (legacy systems are single-user)
7. **DAG-based flexibility** supporting any supply chain topology without custom development
8. **Production-ready** with enterprise features (SSO, RBAC, audit logs) at a fraction of legacy cost

Target Markets

Primary (Direct Replacement Opportunities): - **Manufacturing Companies**: Current Kinaxis/SAP IBP users seeking cost reduction and AI capabilities - **Retail Chains**: Inventory planning teams frustrated with legacy system complexity - **CPG/FMCG Companies**: Multi-tier distribution networks requiring stochastic modeling - **3PL/4PL Providers**: Logistics optimization for multiple clients without per-client licensing - **Mid-Market Enterprises**: Too small for Kinaxis ($5M+ deals) but need enterprise planning

Secondary: - Supply chain consulting firms (replace legacy tools, enable AI consulting) - Business schools (replace expensive academic licenses) - Technology companies (embed AI supply chain planning in existing platforms) - Government/military (replace expensive defense contractor solutions)

Ideal Customer Profile (ICP): - Annual revenue: $100M-$5B - Current planning system: Excel, Kinaxis, SAP IBP, OMP, or homegrown - Pain points: High TCO, slow deployment, poor AI/ML, lack of transparency - Tech maturity: Open to AI but skeptical of "black box" solutions - Decision criteria: ROI, time-to-value, ease of use, vendor independence

---

Success Metrics

Platform Adoption KPIs

- **Active Users**: Monthly active participants, session duration
- **Scenario Volume**: Scenarios run per week, periods per scenario
- **Engagement**: Achievement completion rate, leaderboard participation
- **AI Usage**: % of scenarios with AI opponents, AI suggestion adoption rate

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

- **Training Impact**: Time to proficiency for new hires, knowledge retention scores

Business Impact KPIs

- **Cost Reduction**: Inventory holding costs, backlog penalties
- **Service Level**: Fill rates, stockout frequency, customer satisfaction
- **Efficiency**: Order variability, inventory turnover, cash-to-cash cycle
- **Decision Quality**: Human vs. AI cost differential, improvement over time
- **ROI**: Customer-reported savings vs. platform cost

AI Improvement KPIs

- **Prediction Accuracy**: GNN forecast RMSE, bullwhip prediction error
- **Decision Quality**: AI cost vs. human cost in simulated scenarios
- **Generalization**: Performance on unseen supply chain topologies
- **Explainability**: Human comprehension scores, trust ratings
- **Efficiency**: Training time, inference latency, compute cost

---

Conclusion: Replace Legacy Planning Software with Modern AI

The Continuous Autonomous Planning Platform represents a paradigm shift in how organizations approach supply chain planning:

From Expensive to Affordable: - Legacy systems (Kinaxis/SAP IBP): $6M first-year cost, 18-month deployment - Our platform: $200K first-year cost, 2-4 week deployment (91% TCO reduction)

From Deterministic to Stochastic: - Legacy MRP/DRP: Static calculations with safety stock buffers - Our platform: Full variability modeling (demand, lead time, supplier uncertainty)

From Black Box to Transparent: - Legacy systems: Complex optimization algorithms with poor explainability - Our platform: Humans compete against AI, understand its logic, build trust before deployment

From Rigid to Flexible: - Legacy systems: Months of configuration, expensive custom development - Our platform: DAG-based topology, configure any network in minutes without coding

From Siloed to Collaborative: - Legacy systems: Individual planners optimize in isolation - Our platform: Multi-participant coordination, visibility sharing, joint planning via simulation

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

From Static to Adaptive: - Legacy systems: Periodic model updates by consultants - Our platform: Continuous AI learning from human gameplay and real-world feedback

From Training Burden to Competitive Advantage: - Legacy systems: 3-6 months consultant-led workshops - Our platform: 2-3 weeks gamified learning (60% cost reduction, 3-5x faster adoption)

The Strategic Imperative: Break Free from Legacy Planning Software

In an era of: - **Supply chain volatility** (pandemics, trade wars, climate events requiring stochastic modeling) - **AI adoption pressure** (competitors deploying ML/AI, legacy systems falling behind) - **Cost scrutiny** (CFOs questioning $5M+ planning software investments) - **Skills gaps** (experienced planners retiring, new hires struggling with complex legacy UIs) - **Vendor lock-in fatigue** (proprietary data formats, expensive upgrades, forced migrations)

Organizations need a modern alternative to Kinaxis/SAP IBP/OMP that: 1. **Reduces TCO by 90%** while delivering superior AI/ML capabilities 2. **Deploys in weeks, not years** with pre-built configurations and intuitive UX 3. **Proves AI value** before deployment through risk-free gaming validation 4. **Trains people 60% faster** via simulation vs. consultant-led workshops 5. **Eliminates vendor lock-in** with open architecture and data portability 6. **Models real-world uncertainty** with stochastic variability (demand, lead time, supplier reliability)

The Continuous Autonomous Planning Platform delivers all six—at a fraction of the cost.

Why Now?

Legacy Planning Software Pain Points Reaching Breaking Point: - **Rising Costs**: Kinaxis/SAP license renewals increasing 10-15% annually - **Implementation Failures**: 40-60% of ERP/planning projects fail or exceed budget - **AI Gap**: Legacy vendors offering limited, expensive AI add-ons instead of native capabilities - **Talent Shortage**: Experienced Kinaxis/SAP consultants retiring, new generation prefers modern tools - **Cloud Migration Complexity**: Legacy on-prem to cloud migrations costing millions

Market Opportunity: - Enterprise planning software market: $15B annually - 60% of mid-market companies use Excel (underserved by Kinaxis/SAP pricing) - 30% of Kinaxis/SAP customers exploring alternatives due to cost - AI/ML supply chain optimization market growing 25% CAGR

Competitive Window: - Legacy vendors slow to innovate (18-24 month release cycles) - Startups focused on narrow point solutions (forecasting only, inventory only) - Our platform uniquely combines gaming + AI + planning in one modern architecture

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Next Steps

1. **Deploy Free Tier**: Test with 10-20 users, gather feedback (1 week)
2. **Pilot Program**: Run 3-month pilot with target enterprise customer (Q1 2026)
3. **Case Study**: Document measurable impact, create sales collateral (Q1 2026)
4. **Launch Professional Tier**: Public release with pricing (Q2 2026)
5. **Scale**: 10 enterprise customers, $1M ARR (Q4 2026)
6. **Expand**: Mobile app, 3D visualization, advanced AI (2027)

Appendix: Technical Specifications

System Requirements

Server (Backend + Database): - CPU: 8+ cores (16+ recommended for ML training) - RAM: 16GB minimum (32GB for large scenarios, 64GB for GNN training) - GPU: NVIDIA GPU with 8GB+ VRAM (optional, for GNN training) - Storage: 100GB SSD (database, model checkpoints) - OS: Linux (Ubuntu 22.04 LTS recommended)

Client (Browser): - Modern browser (Chrome 90+, Firefox 88+, Safari 14+, Edge 90+) - 4GB+ RAM - Internet: 5 Mbps+ (for WebSocket real-time updates)

API Endpoints

Scenario Management: - `POST /api/v1/mixed-scenarios/` - Create new scenario - `POST /api/v1/mixed-scenarios/{id}/start` - Start scenario - `POST /api/v1/mixed-scenarios/{id}/execute-period` - Submit decisions, advance period - `GET /api/v1/mixed-scenarios/{id}/state` - Get current scenario state - `GET /api/v1/mixed-scenarios/{id}/history` - Get complete scenario history

AI Agents: - `POST /api/v1/agents/suggest` - Get AI order recommendation - `POST /api/v1/agents/evaluate` - Compare decision against AI - `GET /api/v1/agents/strategies` - List available strategies

Analytics: - `GET /api/v1/analytics/bullwhip` - Calculate bullwhip metrics - `GET /api/v1/analytics/performance` - Get performance report - `POST /api/v1/reports/generate` - Generate comprehensive report

Simulation: - `GET /api/v1/achievements` - List all achievements - `GET /api/v1/participants/{id}/achievements` - Participant achievements - `GET /api/v1/leaderboards` - List leaderboards - `GET /api/v1/leaderboards/{type}/entries` - Leaderboard rankings

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Data Schemas

Scenario State:

```
{
  "game_id": 123,
  "round_number": 15,
  "status": "active",
  "participants": [
    {
      "player_id": 456,
      "role": "Retailer",
      "inventory": 12,
      "backlog": 3,
      "total_cost": 4520,
      "last_order": 8
    }
  ]
}
```

AI Suggestion:

```
{
  "recommended_order": 12,
  "confidence": 0.85,
  "reasoning": "Current inventory (8) + pipeline (10) = 18. Expected demand (14) suggests or
  "alternative_scenarios": [
    {"order": 10, "expected_cost": 520, "risk": "moderate"},
    {"order": 15, "expected_cost": 580, "risk": "low"}
  ]
}
```

Achievement:

```
{
  "achievement_id": "steady_hand",
  "name": "Steady Hand",
  "description": "Maintain order variance below 5 for 15 consecutive rounds",
  "category": "Innovation",
  "rarity": "rare",
  "points": 150,
  "progress": 0.73,
  "unlocked": false
}
```

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

Security & Compliance

Authentication: - JWT tokens with HTTP-only cookies - CSRF protection (double-submit cookie pattern) - MFA support (TOTP via PyOTP)

Authorization: - Role-based access control (SYSTEM_ADMIN, MANAGER, PARTICIPANT) - Resource-level permissions (own scenarios, group scenarios, all scenarios)

Data Protection: - Passwords hashed with bcrypt - Sensitive data encrypted at rest (AES-256) - TLS 1.3 for data in transit - GDPR-compliant data exports and deletion

Audit Logging: - All API requests logged with user, timestamp, IP - Database change tracking (created_at, updated_at, updated_by) - Scenario replay capability for dispute resolution

---

5. Mobile Application: Supply Chain Management On-The-Go

Purpose: Extend platform capabilities to iOS and Android devices for anywhere access to gaming, monitoring, and analytics.

Mobile Capabilities: - **Native iOS & Android**: Single React Native 0.73.2 codebase for both platforms - **Full Feature Parity**: All web features available on mobile (scenarios, templates, analytics, agent monitoring) - **Offline Mode**: Queue decisions and sync when connectivity resumes - **Push Notifications**: Real-time alerts via Firebase Cloud Messaging (scenario events, agent decisions, system updates) - **WebSocket Sync**: Live scenario state updates with Socket.IO integration - **Biometric Auth**: Face ID/Touch ID for secure, fast login - **Mobile-Optimized Charts**: Victory Native charts for real-time analytics - **Agent-to-Agent Monitoring**: View and intervene in AI agent conversations on mobile devices

Technical Architecture: - **Framework**: React Native 0.73.2, React 18.2.0, TypeScript 5.3.3 - **State Management**: Redux Toolkit 2.0.1 with 6 slices (auth, scenarios, templates, analytics, chat, UI) - **Navigation**: React Navigation 6.x (stack + bottom tabs) - **UI Library**: React Native Paper 5.11.6 (Material Design) - **API Client**: Axios 1.6.5 with JWT interceptors - **WebSocket**: Socket.IO Client 4.6.1 for real-time updates - **Offline Storage**: AsyncStorage 1.21.0 with queue-based sync - **Push Notifications**: Firebase Cloud Messaging 19.0.1 - **Charts**: Victory Native 36.9.2 for mobile-optimized visualizations - **Animations**: React Native Reanimated 3.6.1

Key Screens: 1. **Authentication**: Login/register with biometric support 2. **Dashboard**: Overview metrics, active scenarios, quick actions 3. **Scenarios List**: Browse, filter, search scenarios with pull-to-refresh 4. **Scenario Detail**: Real-time scenario state, place orders, period progression 5. **Scenario Detail with A2A Chat**: Monitor agent conversations with human intervention controls 6. **Create Scenario**: Quick Start Wizard or manual configuration 7. **Template Library**: Browse supply chain configurations with

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

previews 8. **Analytics**: Mobile-optimized charts (line, bar, pie), bullwhip metrics, cost breakdown 9. **Profile**: User settings, theme selection, notification preferences

Offline Mode Features: - Queue actions when offline (place orders, create scenarios) - Automatic sync when connectivity resumes - Offline banner indicator - Retry logic with exponential backoff (max 3 retries) - Persistent queue storage via AsyncStorage

Push Notification Types: - **Scenario Events**: "Scenario started", "Period completed", "Your turn", "Scenario finished" - **Agent Events**: "AI agent suggestion", "Agent decision", "A2A message" - **System Events**: "Maintenance scheduled", "Achievement unlocked"

Performance Optimizations: - React.memo for expensive components - FlatList virtualization (initialNumToRender=10, windowSize=5) - Hermes JavaScript engine (Android) - Fast image loading with react-native-fast-image - Lazy loading for large datasets - Native driver animations

Deployment: - **iOS**: App Store deployment via Xcode Archive or Fastlane - **Android**: Google Play Store via AAB bundle - **CI/CD**: Automated builds with GitHub Actions or Fastlane - **OTA Updates**: CodePush for hot updates without app store review

Testing: - Jest 29.7.0 for unit testing - React Native Testing Library 12.4.3 for component testing - 50+ unit tests across slices, screens, and services - Redux Mock Store for isolated state testing

Documentation: Comprehensive 12-file documentation suite (~4,500+ lines): - README.md - Project overview - INSTALL.md - Installation instructions - QUICKSTART.md - 5-minute quick start - QUICK_REFERENCE.md - Command reference - firebase-setup.md - Push notification setup - DEPLOYMENT.md - App Store/Play Store deployment - INTEGRATION_CHECKLIST.md - Pre-launch checklist - TESTING_GUIDE.md - Testing procedures - ACCESSIBILITY.md - Accessibility compliance - A2A_COLLABORATION_GUIDE.md - Agent monitoring features - INDEX.md - Documentation index

Mobile Value Proposition: - **Anywhere Access**: Participate in scenarios and monitor AI agents from mobile devices - **Real-Time Alerts**: Push notifications keep users informed of critical events - **Offline Resilience**: Queue actions offline, sync automatically when connectivity resumes - **Native Performance**: React Native provides 60 FPS native-like experience - **Fast Adoption**: Mobile-first UI reduces training time for field teams - **Lower Barrier to Entry**: No desktop required for supply chain gaming participation

Use Cases: 1. **Field Teams**: Supply chain planners on factory floors can participate in scenarios during downtime 2. **Executive Monitoring**: C-level executives monitor agent decisions on commute 3. **Remote Training**: Employees learn supply chain concepts via mobile gaming 4. **On-Call Intervention**: Emergency response to agent anomalies via

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

push notifications 5. **Mobile Analytics**: Review performance metrics during client meetings

ROI Impact: - **50% faster training adoption**: Mobile simulation increases participation vs. desktop-only - **30% higher engagement**: Push notifications drive 30% more active users - **24/7 monitoring**: Mobile enables round-the-clock agent oversight - **Lower infrastructure costs**: Reduce need for dedicated workstations

---

Deployment & Access

System Access Points

Primary Application URL (Recommended):

```
http://172.29.20.187:8088  (Remote/Network)
http://localhost:8088      (Local)
```

- Main entry point via Nginx proxy
- Handles authentication and routing
- Routes to both frontend and backend API

Mobile Application:

```
iOS App Store: (Pending deployment)
Google Play Store: (Pending deployment)
```

- React Native 0.73.2 iOS and Android apps
- Push notifications via Firebase Cloud Messaging
- Offline mode with automatic sync
- Full feature parity with web platform

API Documentation:

```
http://172.29.20.187:8000/docs  (Swagger UI)
http://172.29.20.187:8000/redoc (ReDoc)
http://localhost:8000/docs      (Local)
```

- Interactive API documentation
- Test endpoints directly from browser
- View request/response schemas

Database Administration (pgAdmin):

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

```
 http://172.29.20.187:5050  (Remote/Network)
http://localhost:5050      (Local)
```

- PostgreSQL database management
- Query interface and schema browser
- Performance monitoring

Default Credentials

System Administrator:

```
 Email:    systemadmin@autonomy.ai
Password: Autonomy@2026
```

- Full system access (SYSTEM_ADMIN role)
- Create/manage groups, users, scenarios
- Access to all admin features
- Model training and configuration

Group Administrators:

```
 TBG Admin:
  Email:    tbg_admin@autonomy.ai
  Password: Autonomy@2026
  Access:   Default TBG group scenarios

Complex SC Admin:
  Email:    complex_sc_admin@autonomy.ai
  Password: Autonomy@2026
  Access:   Complex SC multi-region scenarios

Three FG Admin:
  Email:    ThreeTBG_admin@autonomy.ai
  Password: Autonomy@2026
  Access:   Three finished goods scenarios

Variable TBG Admin:
  Email:    VarTBG_admin@autonomy.ai
  Password: Autonomy@2026
  Access:   Variable demand scenarios
```

Database Admin (pgAdmin):

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

```
 Email:    admin@admin.com
Password: admin
```

Database Connection (PostgreSQL):

```
 Host:     172.29.20.187 (or 'db' from container)
Port:     5432
Database: beer_game
Username: beer_user
Password: change-me-user
```

Quick Start Guide

1. Access the Application:

```
 # Open in browser
http://172.29.20.187:8088

# Or locally
http://localhost:8088
```

2. Login: - Use `systemadmin@autonomy.ai` / `Autonomy@2026` - Navigate to dashboard after login

3. Key Pages:

```
 Main Dashboard:     /dashboard
Scenarios List:     /scenarios
Admin Panel:        /admin
User Management:    /admin/users
Scenario Management: /admin/scenarios
TRM Dashboard:      /admin/trm
Supply Chain Config: /admin/supply-chain-configs
```

4. Run a Scenario: - Navigate to "Scenarios" menu - Select scenario (e.g., "TRM Agent Showcase") - Click "Start Scenario" or "Join Scenario" - Make decisions each round

5. Train TRM Model:

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

```
# Via UI
Navigate to: http://172.29.20.187:8088/admin/trm
Click "Training" tab
Configure and start training

# Via CLI
cd /path/to/backend
python scripts/training/train_trm.py --phase 1 --epochs 10 --device cuda
```

Service Status & Management

Check Running Services:

```
cd /home/trevor/Projects/The_Beer_Game
docker compose ps
```

Service Endpoints: | Service | Port | Status Endpoint | Health |
|---------|------|----------------|--------| | Proxy (Nginx) | 8088 | http://localhost:8088 | Healthy | | Frontend (React) | 3000 | http://localhost:3000 | Healthy | | Backend (FastAPI) | 8000 | http://localhost:8000/health | Healthy | | Database (PostgreSQL) | 5432 | N/A | Healthy | | pgAdmin | 5050 | http://localhost:5050 | Healthy |

Restart Services:

```
# Restart all
make down && make up

# Restart specific service
make restart-backend
make restart-frontend
make proxy-restart

# View logs
make logs
docker compose logs -f backend
docker compose logs -f frontend
```

Network Access

Local Machine (Same Host):

```
http://localhost:8088
```

LAN/Network Access (Other machines on same network):

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

```
http://172.29.20.187:8088
```

- Accessible from any device on local network
- No additional configuration required
- Check firewall if access fails

Internet Access (Public): Requires additional setup: 1. **Port Forwarding**: Forward router port 8088 → 172.29.20.187:8088 2. **Dynamic DNS**: Use No-IP, DuckDNS, or similar 3. **SSL/TLS**: Configure HTTPS with Let's Encrypt 4. **Security**: Enable additional authentication, rate limiting

Or use built-in TLS:

```
make up-tls
# Access via https://172.29.20.187:8443
```

Container Management

Start System:

```
cd /home/trevor/Projects/The_Beer_Game

# Standard startup
make up

# With GPU support
make up FORCE_GPU=1

# With TLS (HTTPS)
make up-tls

# Development mode (hot reload)
make up-dev
```

Stop System:

```
# Stop containers (preserve data)
make down

# Stop and remove volumes (DESTROYS DATA)
docker compose down -v
```

Database Operations:

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

```
# Bootstrap database (first time or after reset)
make db-bootstrap

# Reset scenarios only
make db-reset

# Rebuild database completely
make rebuild-db

# Reseed after rebuild
make reseed-db

# Reset admin password
make reset-admin
```

Troubleshooting Access

Can't Access Application: 1. Check Docker status: `docker compose ps` 2. Check all services are "healthy" 3. Check logs: `make logs` 4. Restart services: `make down && make up`

Port Already in Use:

```
# Check what's using port 8088
sudo lsof -i :8088

# Kill process if needed
sudo kill -9 <PID>

# Or use different port in docker-compose.yml
```

Firewall Issues (Remote Access):

```
# Check firewall status
sudo ufw status

# Allow port 8088
sudo ufw allow 8088/tcp

# Or disable firewall (testing only!)
sudo ufw disable
```

Database Connection Failed:

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

```
# Check database is running
docker compose ps db

# Check database logs
docker compose logs db

# Restart database
docker compose restart db
```

502 Bad Gateway (Nginx): - Backend service not responding - Check backend logs: `docker compose logs backend` - Restart backend: `make restart-backend`

Performance Optimization

For Training (GPU):

```
# Enable GPU mode
make up FORCE_GPU=1

# Verify GPU access
docker compose exec backend nvidia-smi
```

For Production:

```
# Use production compose file
docker compose -f docker-compose.yml -f docker-compose.prod.yml up -d

# Enable Gunicorn workers (in prod config)
# Scale frontend replicas
docker compose up --scale frontend=3 -d
```

Backup & Restore

Backup Database:

```
# Export database
docker compose exec db pg_dump -U beer_user beer_game > backup_$(date +%Y%m%d).sql

# Or use pgAdmin backup tool
# Connect to http://localhost:5050
# Right-click database → Backup
```

Restore Database:

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

```
# Import backup
cat backup_20260117.sql | docker compose exec -T db psql -U beer_user beer_game

# Or use pgAdmin restore tool
```

Monitoring

View Real-Time Logs:

```
# All services
docker compose logs -f

# Specific service
docker compose logs -f backend
docker compose logs -f frontend
docker compose logs -f proxy

# Last 100 lines
docker compose logs --tail=100 backend
```

Resource Usage:

```
# Container stats
docker stats

# Disk usage
docker system df

# Clean up unused resources
docker system prune -a
```

Security Best Practices

Production Checklist: - [ ] Change default passwords (`.env` file) - [ ] Enable HTTPS/TLS (`make up-tls`) - [ ] Configure firewall rules - [ ] Set up rate limiting (Nginx) - [ ] Enable MFA for admin accounts - [ ] Regular database backups - [ ] Update SECRET_KEY in `.env` - [ ] Disable DEBUG mode in production - [ ] Configure CORS properly - [ ] Use strong JWT secrets

Update Passwords:

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.

```
 # Edit .env file
nano /home/trevor/Projects/The_Beer_Game/.env

# Change:
POSTGRESQL_PASSWORD=change-me-user      → your-strong-password
POSTGRES_PASSWORD=change-me-user        → your-strong-password
SECRET_KEY=change-me-secret-key         → your-random-secret-key

# Restart services
make down && make up
```

Contact & Resources

Product Team: - Project Lead: [Name] - AI/ML Lead: [Name] - Frontend Lead: [Name]

Documentation: - User Guide: `/docs/user-guide.md` - API Reference: `/docs/api-reference.md` - Admin Guide: `/docs/admin-guide.md` - Developer Guide: `/docs/developer-guide.md`

Support: - Community Forum: [URL] - Email: support@beergame.ai - Slack Channel: [URL]

Demo: - Live Demo: https://demo.beergame.ai - Login: demo@example.com / Demo@2026 - Video Walkthrough: [YouTube URL]

Document Version: 2.2 **Last Updated**: January 17, 2026 **Status**: Production Ready **Next Review**: April 15, 2026

© 2026 Azirella Ltd. All rights reserved worldwide. Strictly confidential and proprietary.